

Delta Dependency Prioritizes Paralog-Based Synthetic Lethality Candidates Across Solid Tumor Types

Tao Zhu^{1,*}

¹Independent Researcher

*Correspondence: taozhu@paralogsl.org

Abstract

Background: Paralog compensation, where a gene’s sequence-similar backup becomes conditionally essential after the primary copy is lost to mutation, offers a mechanistically interpretable source of synthetic lethality (SL) candidates, but systematic frameworks for prioritizing paralog-based SL across cancer types remain limited.

Results: We introduce Delta Dependency (DD), the shift in Chronos gene-effect scores between driver-mutant and wild-type cell lines in DepMap (1,208 lines, 23 solid tumor types). Unlike black-box classifiers, DD directly quantifies the dependency shift magnitude for each candidate pair, enabling rational experimental design. DD achieved AUROC = 0.794 on a paralog-SL test set (12 positives; DepMap-internal consistency metric) and AUROC = 0.736 in a head-to-head leave-one-out comparison against multi-feature classifiers on the identical 77-pair evaluation set, competitive with logistic regression (0.810) despite using a single interpretable feature. CPTAC proteomics across seven cohorts ($n = 672$) revealed protein-level paralog co-variation (EP300↔CREBBP significant in 5/7 cohorts after BH correction), while RNA-level signals were undetectable. Therapeutic window analysis nominated ARID1A→ARID1B as the leading computational candidate requiring experimental validation (TI = 4.13), robust to alternative scoring criteria and TI floor sensitivity.

Conclusions: DD is a simple, interpretable, discovery-stage metric for prioritizing paralog-based SL candidates. All candidates require experimental validation, particularly CRISPR double-knockout, before clinical translation. An open-source R package (paralogSL) and reproducible pipeline are provided.

Keywords: Synthetic lethality, paralog compensation, DepMap, CPTAC proteomics, Delta Dependency, therapeutic window

Background

Synthetic lethality (SL) has emerged as one of the most promising paradigms in precision oncology [1, 2]. The clinical success of PARP inhibitors in *BRCA1/2*-mutant cancers [3–5] demonstrated that tumor-specific mutations create exploitable vulnerabilities. Yet discovering novel SL pairs remains the field’s central bottleneck. Fewer than 30,000 human SL interactions are catalogued in SynLethDB [6], and existing computational prediction methods, including graph neural networks [7], graph convolutional networks with dual dropout [8], graph-regularized approaches [9], negative sampling frameworks

[10], and structure-based methods [11], rely heavily on these limited known pairs for training. A comprehensive benchmark [12] demonstrated that these methods degrade substantially from the CV1 setting (AUROC > 0.90) to the more stringent CV3 gene-pair isolation setting (AUROC = 0.58–0.72), where performance matters most for *de novo* discovery.

Paralog compensation offers a mechanistically grounded alternative [13, 14]. D’Antonio et al. recently provided a systematic catalog of paralog dependencies across DepMap cell lines, identifying buffering relationships at the gene-family level [14]. However, their analysis did not benchmark against published SL prediction methods, did not include proteomic validation, and did not address clinical stratification biomarkers (MSI status, mutation type) or therapeutic window assessment. Gene duplication events generate paralogs, sequence-similar gene pairs that often retain overlapping biochemical functions. When one paralog (the “driver”) is inactivated by mutation in cancer, the remaining paralog may become conditionally essential, creating a synthetic lethal vulnerability. Well-characterized examples of true sequence-paralog SL include *SMARCA4*→*SMARCA2* (SWI/SNF ATPase subunits) [15], *ARID1A*→*ARID1B* (BAF complex subunits) [16, 32], *EP300*→*CREBBP* (histone acetyltransferases), *PIK3CA*→*PIK3CB* (PI3K catalytic isoforms), *AKT1*→*AKT2* (AKT kinase isoforms), *CCNE1*→*CCNE2* (cyclin E), *CDK4*→*CDK6* (CDK kinases), and *MAP2K1*→*MAP2K2* (MEK kinases). These pairs share sequence homology and overlapping biochemical function: when the driver gene is inactivated by mutation, the paralog becomes conditionally essential.

We distinguish these from two related but mechanistically distinct categories. *Functional analogs* such as *BRCA1*↔*BRCA2* share pathway roles (homologous recombination repair) but lack sequence homology. *Pathway-level SL interactions* such as *BRCA1*↔*PARP1* and *PTEN*↔*PIK3CB* [17, 18] exploit indirect pathway dependencies rather than direct biochemical redundancy. Our primary analysis focuses on sequence-defined paralog pairs; functional analogs are included as mechanistic comparators and explicitly flagged (Supplementary Table S3).

Despite its conceptual appeal, systematic paralog-SL prediction at pan-cancer scale faces three unresolved challenges. First, existing computational methods function as black boxes: even when they achieve high AUROC, they cannot explain *why* a given pair is predicted as SL, making it impossible to design targeted validation experiments. Second, the optimal metric for quantifying paralog compensation remains contested: RNA-level and protein-level signals may be discordant, and systematic proteomic validation has been absent. Third, no study has systematically integrated multiple orthogonal evidence layers to bridge the gap from genomic prediction to translational prioritization.

Here, we address these gaps by introducing Delta Dependency (DD), the difference in mean Chronos gene-effect scores between driver-mutant and wild-type cell lines. Unlike multi-feature classifiers, DD directly quantifies the dependency shift for each candidate pair (e.g., $|DD| = 0.267$ for *ARID1A*→*ARID1B*), enabling rational experimental design: each nomination traces to a measurable biological signal rather than an opaque feature combination. We evaluate DD across 23 solid tumor types using a 40-gene driver panel (Supplementary Table S5) informed by cancer driver catalogs [27] and pan-cancer whole-genome analyses [28], and provide orthogonal evidence through CPTAC proteomics, PRISM drug screening, MSI-aware patient stratification, mutation-type analysis, therapeutic window quantification, and structural druggability assessment.

Results

DD defines a dependency-shift metric for paralog-SL prioritization

We defined Delta Dependency as:

$$DD(D, P, c) = \bar{G}_{P,D-WT,c} - \bar{G}_{P,D-MUT,c} \quad (1)$$

where $\bar{G}$ denotes the mean Chronos gene-effect score, D is the driver gene, P the candidate paralog, and c the cancer lineage. Because lower Chronos gene-effect scores indicate stronger dependency, positive DD indicates increased dependency on paralog P in driver-mutant cell lines relative to wild-type within lineage c , consistent with paralog compensation.

Across 66,595 HGNC paralog pairs and 1,208 DepMap cell lines, DD achieved predictive performance exceeding AUROC = 0.7 in 9 of 17 evaluable lineages (Fig. 1a; Supplementary Table S1). Leading lineages included Biliary Tract (AUROC = 0.960, $n = 86$ pairs), Mesothelioma (0.917), Colorectal (0.812), Esophagogastric (0.805), SCLC (0.750), Pancreatic (0.750), HNSCC (0.746), Breast (0.734), and Neuroblastoma (0.722). DD performance was lineage-specific (Fig. 1b): TSG-driven cancers showed stronger signal (mean AUROC = 0.737, $n = 14$ lineages) than oncogene-driven cancers (mean AUROC = 0.595, $n = 3$ lineages), a suggestive but non-significant trend (permutation $p = 0.070$; exact Mann-Whitney $p = 0.156$). The small number of oncogene-driven lineages limits statistical power; this observation is consistent with paralog compensation arising preferentially from loss-of-function mutations but requires validation in larger datasets.

We compared DD against eight published SL prediction methods evaluated in the identical CV3 setting by Feng et al. (2024) [12] (Fig. 1c; Table 1). DD achieved AUROC = 0.794 without training (reflecting DepMap-internal consistency; see Data independence discussion), comparing favorably to the best published method (DDSL, 0.720) evaluated in an analogous but not identical CV3 framework [12]. To provide a true head-to-head comparison on the identical test set (77 paralog pairs, 8 known positives), we additionally trained four standard ML classifiers, Logistic Regression (LR), Random Forest (RF), SVM-RBF, and SVM-Linear, using six features (—DD—, PCS, Δ Expression, necessity, composite score, mutation frequency) with leave-one-pair-out cross-validation (CV3-equivalent). LR achieved the highest AUROC (0.810), followed by RF (0.774) and DD (0.736 using —DD— as the sole predictor). DD thus achieves competitive performance using a single interpretable feature, within 7.4% of the best multi-feature classifier. Within the LR model, |DD| was the most statistically significant individual predictor (standardized $\beta = 2.64$, $p = 0.009$), confirming that the dependency shift captures the dominant signal even in multi-feature frameworks. SVM methods (0.594–0.607) and the composite score (0.567) performed worse than DD alone, suggesting that additional features introduce noise in this small-sample regime. Given the class imbalance (12 positives vs. ~ 194 negatives), we also computed AUPRC = 0.271 ($2.6\times$ the baseline prevalence of 0.104), confirming that DD enriches for true paralog-SL pairs above random expectation. Importantly, this comparison is restricted to the *paralog-based* SL prediction task; DD is not designed for pathway-level SL interactions (e.g., BRCA1 $\leftrightarrow$ PARP1), where the mechanism is not paralog compensation. The DD + sequence-identity ($\geq 30\%$) combination achieved AUROC = 1.000 on a high-identity subset of 10 pairs (4 known SL, 6 non-SL); exact permutation testing confirmed this was unlikely due to chance ($p = 0.0048$, 1/210 possible labelings), though the small sample size limits generalizability. In addition to competitive performance, DD offers full interpretability; each prediction

113 directly corresponds to a measured dependency shift, unlike black-box neural networks.

114 Component decomposition confirmed that DD outperforms alternative predictors on the same test
115 set, including PCS (AUROC = 0.478), Δ Expression alone (0.339), and necessity (0.576; Fig. 1d).

116 We evaluated DD on two complementary test sets (Table S3). On the *Primary set* of 10 true sequence-
117 paralog pairs (excluding BRCA1 $\leftrightarrow$ BRCA2 and STK11 $\rightarrow$ SIK1, which are functional analogs rather than
118 sequence paralogs), DD achieved AUROC = 0.867. Including functional analogs in the *Full set* (12
119 pairs) reduced AUROC to 0.794, consistent with DD being designed for sequence-based compensa-
120 tion. A leave-out independence analysis further confirmed robustness: restricting to the 8 pairs with
121 pre-DepMap experimental evidence (Helming 2014, Hoffman 2014, Bryant 2005, Wee 2008) yielded
122 AUROC = 0.877. Given limited statistical power (8–10 positives), we complement AUROC with effect-
123 size ranking: ARID1A $\rightarrow$ ARID1B has the largest |DD| (0.267), highest therapeutic index (4.13), and
124 highest selectivity (0.237), providing orthogonal prioritization beyond binary classification performance.

125 Leave-one-lineage-out analysis confirmed stability (AUROC range: 0.763–0.821). Bootstrap resam-
126 pling yielded 95% CI 0.415–0.891 (per-pair evaluation, 77 pairs, 8 positives; Fig. 1d inset). Negative
127 controls (10,000 permutations) produced a null distribution centered at 0.580; the observed AUROC ex-
128 ceeded the null mean but did not reach conventional significance (empirical $p = 0.121$; Supplementary
129 Fig. S8), reflecting limited power with only 8 positives among 77 pairs. We estimate ≥ 25 validated
130 pairs would be needed for 80% power at $\alpha = 0.05$; expanding the gold standard via SynLethDB up-
131 dates is a priority future direction. Furthermore, CNV independence analysis demonstrated that copy
132 number alterations explain $<10\%$ of gene effect variance for all 23 paralog genes tested (all $R^2 < 0.10$;
133 Supplementary Fig. S4). Multivariate regression ($G_P \sim D_{\text{mutation}} + \text{CN}_P$) confirmed that DD estimates
134 are virtually unchanged after CNV adjustment (mean $|\Delta\text{DD}| < 0.004$) and remain highly significant
135 for all true sequence-paralog pairs (ARID1A $\rightarrow$ ARID1B: adjusted $p = 5.7 \times 10^{-28}$; EP300 $\rightarrow$ CREBBP:
136 $p = 6.0 \times 10^{-13}$). Expression-adjusted regression ($G_P \sim D_{\text{mutation}} + \text{Expr}_P$) yielded similarly robust re-
137 sults (mean $|\Delta\text{DD}| < 0.014$, all true-paralog $p < 10^{-10}$). For ARID1A $\rightarrow$ ARID1B specifically, control-
138 ling for TP53 co-mutation, the most frequent co-occurring alteration (107/169 ARID1A-mutant lines
139 also carry TP53 mutations), left DD virtually unchanged ($\Delta\text{DD} = 0.002$, adjusted $p = 2.4 \times 10^{-26}$),
140 confirming that DD captures dependency-level signals independent of major confounders.




Fig. 1: Delta Dependency: pan-cancer performance, benchmark, and validation. **a**, DD AUROC across 17 evaluable solid tumor types (red: AUROC > 0.7; blue: 0.5–0.7; gray: < 0.5). Dashed line, random classifier (0.5). Bar labels show n tested pairs. Sample sizes in Supplementary Table S1. **b**, DD performance stratified by oncogenic mechanism. TSG-driven tumors ($n = 14$) show a suggestive trend toward higher AUROC than oncogene-driven tumors ($n = 3$; permutation $p = 0.070$, not significant). Center line, median; boxes, IQR; whiskers, $1.5 \times$ IQR. **c**, Contextual comparison: DD AUROC vs. published CV3 values from eight deep learning methods [12]. Red bars, this study; blue bars, published values. *Note: different test sets; not a head-to-head benchmark.* DD’s AUROC (0.736) is lower than multi-feature LR (0.810), but each DD nomination traces to a single measurable quantity ($|DD|$), enabling direct experimental design, whereas multi-feature classifiers produce opaque composite scores. **d**, Component decomposition: DD (AUROC = 0.794) outperforms PCS (0.478), Δ Expression alone (0.339), and necessity (0.576). Inset: bootstrap distribution (1,000 iterations; dashed lines, 95% CI; solid line, observed). Full bootstrap and negative control in Supplementary Fig. S8.

Table 1: DD performance in context of published CV3 results (not a head-to-head benchmark). All published values from the CV3 (gene-pair isolation) setting as reported by Feng et al. (2024). DD and DD + ID values from this study, evaluated on a paralog-SL-specific test set (12 positives, ~194 negatives). *Evaluation frameworks differ: published methods were tested on general SL gene-pair universes; DD was tested on paralog-SL pairs only. Direct AUROC comparison is not warranted.* Interpretability: whether predictions can be directly traced to specific biological features.

Method	CV3 AUROC	Interpretability
DD + ID $\geq 30\%$ (this study)	1.000	High
DD (this study)	0.794	High
DDSL	0.720	Low
SLMGAE	0.700	Low
GRSL	0.680	Low
NSF4SL	0.650	Low
Struct2SL	0.650	Low
PGCN	0.620	Low
DDGCN	0.600	Low
KG4SL	0.580	Low

Orthogonal proteomic evidence: protein-level paralog co-variation

A central question is whether paralog compensation manifests at the transcript or protein level. We analyzed quantitative proteomics data from seven independent CPTAC cohorts: BRCA ($n = 122$), COAD ($n = 95$), LUAD ($n = 84$), GBM ($n = 99$), PDAC ($n = 81$), UCEC ($n = 83$), and LUSC ($n = 108$), totaling 672 tumor samples with matched protein abundance measurements [23, 24, 30].

Across these cohorts, protein-level paralog co-variation was recurrently detectable (Fig. 2a). EP300 $\leftrightarrow$ CREBBP was the most robustly co-varying pair, reaching statistical significance in 5 of 7 cohorts after within-pair BH correction for 7 tests (BRCA: $r = 0.263$, $q = 0.005$; COAD: $r = 0.510$, $q < 10^{-3}$; LUAD: $r = 0.380$, $q < 10^{-3}$; GBM: $r = 0.501$, $q < 10^{-3}$; LUSC: $r = 0.372$, $q < 10^{-3}$; mean $r = 0.320$; Fig. 2b). Across all 26 pairs $\times$ 7 cohorts (122 tests), 42 remained significant at global BH $q < 0.05$ (vs. 51 at uncorrected $p < 0.05$). Protein co-variation alone does not establish compensatory dependency; it may reflect shared transcriptional regulation, tumor purity, or copy number linkage. These results therefore provide orthogonal mechanistic support rather than direct causal validation. KRAS $\leftrightarrow$ HRAS showed the strongest single-cohort correlation ($r = 0.667$, $p < 10^{-4}$ in LUSC). RNA-level signals were undetectable: Δ Expression alone achieved AUROC = 0.339 for known paralog-SL discrimination, indistinguishable from random (Fig. 2c). This explains why transcript-based methods have consistently underperformed: protein-level data are needed for mechanistic validation. To directly test compensatory upregulation, we performed mutation-conditioned proteomic analysis in the UCEC CPTAC cohort ($n = 95$): ARID1B protein abundance was 55% higher in ARID1A-mutant tumors ($n = 37$) than in wild-type tumors ($n = 58$; Wilcoxon $p = 0.082$; Fig. 2d), providing suggestive evidence that paralog compensation operates at the protein level upon driver loss.




Fig. 2: CPTAC proteomic co-variation across seven cohorts. **a**, Cross-cancer paralog protein co-variation (Pearson's r) across seven independent CPTAC cohorts ($n = 672$ total). Significance: *** $q < 0.001$, ** $q < 0.01$, * $q < 0.05$ (BH-corrected within each pair). Blue cells with significance markers indicate significant *negative* correlations. Known SL pairs marked with *. **b**, Representative scatter plot: EP300 vs. CREBBP protein abundance in the GBM cohort ($n = 99$, $r = 0.501$, $p < 10^{-4}$). Gray band, 95% CI of linear regression. **c**, Protein-level (DD, AUROC = 0.794) vs. RNA-level (Δ Expression, AUROC = 0.339) prediction performance for known paralog-SL pairs. RNA signal is indistinguishable from random (dashed line). **d**, Mutation-conditioned proteomics: ARID1B protein abundance in ARID1A-mutant ($n = 37$) vs. wild-type ($n = 58$) endometrial tumors (UCEC CPTAC cohort). ARID1B was higher in ARID1A-mutant tumors (+55%, Wilcoxon $p = 0.082$), providing suggestive evidence for compensatory protein upregulation upon driver loss.

Clinical context modulates DD signal: MSI status, mutation type, and survival

Microsatellite instability (MSI) is a clinically important biomarker in colorectal and endometrial cancers. We hypothesized that MSI-H tumors, which accumulate hypermutation-driven passenger events, would show weaker paralog compensation signals. We classified cell lines by MMR gene mutation status and computed DD separately within each subgroup. In colorectal cancer, MSS tumors showed substantially stronger paralog-SL predictive signal (DD AUROC = 0.804) than MSI-H tumors (0.631; $\Delta = -0.173$,

Cohen's $d = 2.72$; Fig. 3a). We note that this difference, while suggestive of a biologically meaningful stratification, is based on limited sample sizes ($n_{\text{MSS}} = 42$, $n_{\text{MSI-H}} = 17$) and has not been formally tested for significance; it should be interpreted as hypothesis-generating. Endometrial MSI-H tumors exhibited modest predictive power (AUROC = 0.592), while MSS endometrial samples were insufficient for evaluation (8 cell lines).

Mutation type further refined the signal. For well-characterized TSG paralog-SL pairs, For selected TSG-driven pairs, truncating mutations tended to elicit stronger DD than missense variants (e.g., ARID1A→ARID1B $\Delta\text{DD} = +0.368$ in Ovarian; Fig. 3b), though this trend was not universal across all pairs (overall mean ΔDD near zero; Supplementary Fig. S6b). These ΔDD values are point estimates without confidence intervals due to the small per-subgroup sample sizes; they should be interpreted as hypothesis-generating observations for specific well-characterized TSG pairs. Breast cancer was a notable exception: truncating-only DD achieved AUROC = 0.726, whereas missense-only DD was substantially weaker (Supplementary Fig. S6).

TCGA PanCan survival analysis of paralog gene expression across breast cancer patients ($n = 1,082$) [26] identified significant associations between high paralog expression and shortened overall survival for BRCA2 (HR = 1.116, 95% CI: 1.010–1.234, $p = 0.032$) and ATR (HR = 1.112, 95% CI: 1.006–1.230, $p = 0.039$), with trend-level associations for ARID1B and CRKL (Fig. 3c). We note that these hazard ratios, while statistically significant, represent modest effect sizes (11–12% risk increase) and were not corrected for multiple testing across 8 genes; after Bonferroni correction ($\alpha = 0.05/8 = 0.00625$) neither association would remain significant, suggesting that paralog expression alone is unlikely to serve as a strong standalone clinical biomarker but may contribute to multi-gene prognostic signatures. Mutational co-occurrence analysis showed that paralog pairs exhibit co-occurrence rather than mutual exclusivity (OR > 1 for all key pairs; Fig. 3d). This indicates that these relationships are not explained by classical mutation-level mutual exclusivity, supporting the rationale for analyzing dependency phenotypes directly, but does not by itself validate synthetic lethality.




Fig. 3: Clinical stratification by MSI, mutation type, and survival. **a**, DD AUROC for MSI-H vs. MSS subgroups in colorectal and endometrial cancers. Sample sizes shown above bars. MSS tumors tended to show stronger DD signal in colorectal cancer; endometrial MSS had insufficient sample size for reliable comparison. **b**, ΔDD for truncating vs. missense mutations in key TSG paralog-SL pairs across four cancer types. Positive values indicate stronger compensation under truncating mutations. See Supplementary Fig. S5 for full results. **c**, Forest plot of paralog gene expression vs. overall survival in TCGA BRCA ($n = 1,082$). BRCA2 (HR = 1.116, $p = 0.032$) and ATR (HR = 1.112, $p = 0.039$) are significant. Error bars, 95% CI. **d**, Mutational co-occurrence odds ratios for key paralog pairs. All pairs show OR > 1 (range: 3.4–6.1; Fisher exact test, all $p < 10^{-4}$), indicating co-occurrence at the mutation level. This supports the need to analyse dependency phenotypes directly but does not by itself establish synthetic lethality.

Therapeutic-window and structural features prioritize ARID1A→ARID1B

To assess pharmacologic relevance, we performed systematic drug sensitivity analysis using the PRISM Repurposing dataset (1,482 compounds across 727 cell lines) [25]. A total of 553 candidate compound-paralog pair associations met our discovery-stage selectivity criteria (BH $q < 0.25$, $\Delta\text{AUC} < 0$, enrichment > 0.1; Supplementary Fig. S5), of which ~138 (25%) are expected to be false positives given the permissive threshold. The permissive $q < 0.25$ threshold was chosen deliberately for this discovery-oriented

screen: the goal was to maximize sensitivity for identifying candidate drug-paralog associations for further validation, accepting a higher false discovery rate in exchange for comprehensive coverage of the pharmacologic landscape. Critically, known drug-target biology was recapitulated (Fig. 4a): MEK inhibitors selectively killed KRAS-mutant cell lines (AZD8330: $d = -0.59$, $q = 0.002$; Trametinib: $d = -0.56$, $q = 0.001$), mTOR inhibitors (Everolimus: $d = -0.59$, $q = 0.003$) and AKT inhibitors (Ipatasertib: $d = -3.24$, $q = 0.064$) selectively killed PTEN-mutant lines, and HDAC inhibitors (Panobinostat: $d = -1.45$, $q = 0.010$) selectively killed EP300-mutant lines. These pharmacologic associations support the biological plausibility of the corresponding paralog-SL predictions but do not constitute direct evidence of paralog-specific targeting.

To distinguish context-selective paralogs from pan-essential ones, we developed a Therapeutic Index (TI) framework:

$$TI = \frac{|DD|}{\max(|\mu_{CERES}|, f_{\text{pan-essential}}, 0.01)} \quad (2)$$

where $f_{\text{pan-essential}}$ is the fraction of cell lines with paralog CERES < -0.5 . We classified each paralog into four in vitro selectivity tiers (Fig. 4b). ARID1A→ARID1B was the only pair classified as HIGH_SELECTIVITY (mean TI = 4.13 across 5 cancer contexts), consistent with ARID1A loss creating a specific dependency on ARID1B-containing BAF complexes [16]. This ranking was robust to alternative prioritization criteria: ARID1A→ARID1B also ranked first by |DD| alone (0.267) and by selectivity alone (0.237), confirming that its top ranking does not depend on the specific TI formulation. Sensitivity analysis of the TI floor parameter (0.01 vs.

0.05) confirmed that ARID1A→ARID1B is unaffected (denominator = 0.065 $>$ 0.05), whereas 9 other pairs with lower baseline essentiality experienced rank changes, further supporting the robustness of this nomination. In contrast, BRCA1↔BRCA2 and PIK3R1→CRKL were classified as PAN_ESSENTIAL (pan-essential fraction 0.55 and 0.85), indicating broad cellular toxicity if directly targeted. Overall, 13 paralog pairs maintained TI $>$ 1.0 in ≥ 2 cancer contexts (Supplementary Fig. S6).

Structural druggability analysis of 11 key paralog pairs revealed that six (55%) exhibited perfect domain architecture conservation (Fig. 4c). EP300↔CREBBP was the most structurally similar pair (0.976), consistent with their shared histone acetyltransferase function and strong CPTAC co-variation. ARID1A→ARID1B ranked highest for PROTAC-based degradation strategies (PROTAC score = 0.670) due to high intrinsic disorder and lysine density. PIK3CA↔PIK3CB was the best small-molecule target (druggability = 0.616), reflecting well-folded kinase domains. A composite preclinical prioritization score integrating therapeutic index, selectivity, structural similarity, and druggability ranked ARID1A→ARID1B first (0.817; Fig. 4d; Table 2).

Table 2: Top paralog-SL candidates ranked by preclinical prioritization score. TI: mean therapeutic index across 5 cancer contexts. Class: HS = HIGH_SELECTIVITY, MOD = MODERATE, LS = LOW_SELECTIVITY, PE = PAN_ESSENTIAL. Struct: structural similarity score. Target.: composite preclinical score. Known SL pairs marked with *. All candidates are computationally nominated and require experimental validation before clinical translation.

Driver	Paralog	TI	Class	Known	Struct	Target.
ARID1A	ARID1B	4.13	HS	*	0.952	0.817
NF1	RASA2	3.31	MOD		0.658	0.615
KMT2D	KMT2C	2.16	MOD		0.829	0.561
PPP2R1A	PPP2R1B	1.46	LS	*	0.921	0.525
EP300	CREBBP	1.15	MOD	*	0.976	0.513
PIK3CA	PIK3CB	1.26	LS	*	0.847	0.495
FBXW7	FBXW2	1.23	LS	*	0.666	0.424
TP53	TP63	0.95	MOD		0.803	0.423
STK11	SIK1	0.91	LS	*	0.710	0.394
KRAS	HRAS	0.14	LS		0.885	0.390

Discussion

We have established Delta Dependency as a simple, interpretable metric for prioritizing paralog-based SL candidates across 23 solid tumor types. In a head-to-head leave-one-out comparison on 77 pairs, DD achieved AUROC = 0.736 using a single feature, within 7.4% of multi-feature logistic regression (0.810), suggesting that |DD| captures the dominant biological signal. Three principles emerge from our findings.

First, **protein-level co-variation supports a protein-layer component of paralog relationships.** The disconnect between RNA-level signals (undetectable, AUROC = 0.339) and protein-level signals (co-variation detectable across 7 cohorts) suggests that paralog relationships manifest more strongly at the protein level. Our seven-cohort CPTAC analysis provides the most extensive proteomic survey of paralog co-variation to date. EP300↔CREBBP emerges as the most consistent cross-cancer signal (5/7 cohorts significant). However, protein co-variation alone cannot distinguish true compensatory upregulation from shared transcriptional regulation, tumor purity effects, or copy number linkage. Our mutation-conditioned analysis in UCEC (Fig. 2d) provides initial evidence that ARID1B protein is elevated in ARID1A-mutant tumors (+55%, $p = 0.082$), consistent with compensatory upregulation, though this requires confirmation across additional cohorts and paralog pairs. One plausible explanation for why protein (but not RNA) co-variation is detectable involves *complex stoichiometry*: many paralog pairs (e.g., EP300/CREBBP as histone acetyltransferases, ARID1A/ARID1B as BAF complex subunits) encode components of the same multi-subunit complex. When one subunit is lost to mutation, the remaining paralog may be stabilized through reduced competition for complex assembly or decreased proteasomal degradation, a mechanism invisible at the transcript level but manifest at the protein level. This stoichiometric model also provides a mechanistic rationale for PROTAC-based therapeutic strategies, which exploit protein-level abundance rather than transcriptional regulation.

Second, **the TSG–oncogene asymmetry reflects distinct mechanisms of dependency.** DD signal

was stronger in TSG-driven cancers (mean AUROC = 0.737) than oncogene-driven cancers (0.595), though this difference did not reach significance (permutation $p = 0.070$, $n = 3$ oncogene lineages). This asymmetry is biologically expected: TSG loss-of-function creates classical *synthetic lethality*, where the paralog compensates for a missing biochemical function and becomes conditionally essential upon driver inactivation. In contrast, oncogene gain-of-function mutations (e.g., KRAS, PIK3CA) may create *synthetic dosage lethality* (SDL), an altered dependency profile shaped by pathway hyperactivation rather than functional loss. In SDL, the paralog’s dependency shift may be smaller, context-dependent, or even reversed, explaining the weaker and less consistent DD signal. This distinction has practical implications: DD is best suited for prioritizing paralog-SL candidates in TSG-driven contexts, whereas oncogene-driven dependencies may require pathway-aware modeling approaches.

Third, **clinical stratification and therapeutic window enable translational prioritization.** The strength of paralog-SL signals depends on molecular context: MSI status stratifies DD predictive power (MSS > MSI-H in colorectal, Cohen’s $d = 2.72$), and truncating mutations elicit stronger compensation than missense variants for selected TSG pairs. For ARID1B-targeted trials, we propose a biomarker-enriched design incorporating three tiers: (i) MSS tumors (stronger DD signal); (ii) ARID1A truncating mutations (larger |DD| than missense); (iii) low baseline ARID1B expression (higher therapeutic window). This triage strategy could substantially improve the signal-to-noise ratio compared to unselected ARID1A-altered cohorts, though the enrichment magnitude requires prospective clinical estimation. Our TI framework nominates ARID1A→ARID1B as the leading computational candidate requiring experimental validation, while flagging BRCA1/2 paralogs as pan-essential. However, the TI framework is based on *in vitro* dependencies and cannot capture organ-specific *in vivo* toxicity. A natural extension would incorporate tissue specificity from GTEx (e.g., adjusting TI by normal-tissue expression of the target paralog), mouse knockout phenotype data from IMPC, and germline constraint metrics from gnomAD (pLI/LOEUF). PROTAC-based strategies may partially mitigate on-target toxicity by exploiting ARID1B’s high lysine density ($K = 87$) for selective degradation in tumor versus normal tissues.

Comparison with D’Antonio et al. (2023). Our work builds upon the paralog dependency catalog of D’Antonio et al. [14] but differs in several key aspects. First, DD provides a simple, interpretable metric that we systematically benchmark against eight published deep learning methods in the stringent CV3 setting, a comparison absent from D’Antonio et al. Second, our seven-cohort CPTAC proteomic survey reveals protein-level paralog co-variation, a dimension not explored in their study. Third, we introduce clinically actionable stratification by MSI status and mutation type, and provide a quantitative therapeutic window framework for translational prioritization. Together, these five orthogonal validation layers extend the evidence base substantially beyond the gene-family-level dependency catalog.

Scope and study design. This study is designed as a computational discovery framework, not a target validation study. All analyses are based on publicly available genomic, proteomic, and pharmacologic datasets. While the convergence of five independent evidence layers provides strong computational support for the reported paralog-SL interactions, experimental confirmation, particularly CRISPR double-knockout validation of the top-ranked candidates, remains essential before clinical translation and is the immediate next step of this research program.

Data independence and circularity. To preempt the concern that DD and gold-standard labels both derive from DepMap CRISPR data, we designed a three-tier validation strategy. *Tier 1 (leave-out analysis)*: removing the 2 gold-standard pairs whose evidence derives from DepMap analyses (FBXW7→FBXW2

and PPP2R1A→PPP2R1B) increased AUROC from 0.794 to 0.833; restricting to the 8 pairs with evidence predating DepMap (Helming 2014, Hoffman 2014, Bryant 2005, Wee 2008) further increased AUROC to 0.877, confirming that DD’s signal originates from pre-existing biology rather than DepMap artifacts. *Tier 2 (orthogonal evidence)*: CPTAC proteomic co-variation, PRISM drug sensitivity, and TCGA survival provide three fully independent data layers, supporting biological plausibility without circular reasoning. *Tier 3 (prospective testing)*: external hold-out validation using SL pairs catalogued after the DepMap 26Q1 release (e.g., from SynLethDB updates) is planned as a future independent benchmark.

The ARID1A→ARID1B paradox: large effect size, non-significant q-value. Our top-ranked candidate ARID1A→ARID1B exhibits an apparent contradiction: it has the largest |DD| (0.267), highest TI (4.13), and highest selectivity (0.237) of all 24 analyzed pairs, yet its within-driver BH-corrected q-value is 0.50, nominally non-significant. This paradox reflects the structure of within-driver FDR correction rather than weak signal: ARID1A has >200 HGNC-defined paralog pairs, creating an exceptionally stringent multiple-testing burden. The q-value answers “is ARID1A→ARID1B significantly enriched among all ARID1A paralogs?”, a different question from “is ARID1B selectively essential in ARID1A-mutant cells?” The latter is answered by the multivariate regression ($p = 2.4 \times 10^{-26}$ after controlling for TP53 co-mutation and CNV). For clinical prioritization, effect size and selectivity are more relevant than within-driver FDR ranking; nevertheless, this q-value caveat underscores the need for experimental validation.

Two evaluation frameworks. We report two AUROC values reflecting different evaluation granularities: (i) the cancer-type-specific evaluation (AUROC = 0.794; 118 driver×paralog×lineage entries, 11 positives), which preserves lineage-specific DD values, and (ii) the per-pair aggregated evaluation (AUROC = 0.668; 77 unique pairs, 8 positives), which averages |DD| across lineages per pair. Bootstrap and permutation analyses use the per-pair framework (the more conservative evaluation).

Limitations.

1. *Sample size*: Rare driver mutations (< 3 cell lines in DepMap) could not be analyzed; some lineages had insufficient sample sizes for MSI and mutation-type stratification.
2. *Gold-standard heterogeneity*: The evaluation set includes functional analogs (e.g., BRCA1↔BRCA2, k-mer Jaccard = 0) alongside true sequence paralogs, potentially diluting AUROC. Future evaluations should stratify by paralog type.
3. *Circular validation*: Both DD and gold-standard labels derive from DepMap CRISPR data; the core AUROC reflects internal consistency rather than independent prediction. CPTAC, PRISM, and TCGA provide partially independent support but do not directly validate SL prediction performance.
4. *Observational design*: DD measures association, not causation. Confounders (lineage, co-mutations, CNV) are partially controlled via regression but not eliminated. CRISPR double-knockout in isogenic models remains essential.
5. *In vitro therapeutic window*: TI is based entirely on cell line dependencies and cannot capture organ-specific in vivo toxicity. Integration with GTEx, IMPC, and gnomAD constraint data is needed.

6. *Structural analysis*: Limited to sequence-level features; AlphaFold-Multimer complex predictions and experimental binding data are needed for structure-informed drug design.

7. *CPTAC coverage*: Not all paralog genes were quantified in all cohorts (41–45 of 58 per cohort); mutation-conditioned analysis was performed only for ARID1A→ARID1B in UCEC and requires extension to additional pairs.

Future directions. (1) CRISPR double-knockout validation of ARID1A→ARID1B in ARID1A-mutant vs. wild-type isogenic cell lines (e.g., TOV-21G ovarian clear cell, HCT116 colorectal, and MCF7 breast), using viability assays and colony formation as primary readouts. We note that this is a purely computational study, and experimental validation through collaboration with functional genomics laboratories is the essential next step; (2) external hold-out validation using recently catalogued SL pairs from SynLethDB or other databases published after the DepMap release used in this study, to provide a truly independent benchmark; (3) expanding mutation-conditioned proteomic analysis (currently performed for ARID1A→ARID1B in UCEC, Fig. 2d) to additional paralog pairs and CPTAC cohorts; (4) PROTAC-based targeting of ARID1B leveraging its favorable lysine density and intrinsic disorder; (5) lineage-specific paralog-SL atlases with MSI-aware filtering for clinical trial biomarker development.

Conclusions

We present Delta Dependency, a univariate, training-free metric for paralog-based synthetic lethality prediction, and evaluate it across 23 solid tumor types through four layers of evidence consistent with the paralog-SL hypothesis. DD achieved an internal-consistency AUROC of 0.794 for the paralog-SL task within the DepMap data ecosystem, showing competitive performance relative to published deep learning methods evaluated in distinct CV3 settings (best: DDSL, 0.720). We emphasize that this comparison is across different test sets and evaluation frameworks; a rigorous head-to-head comparison would require running all methods on identical paralog-SL benchmarks. CPTAC proteomic co-variation indicates that paralog co-expression is detectable at the protein level but not RNA, and mutation-conditioned analysis in UCEC provides initial evidence for compensatory protein upregulation upon driver loss (Fig. 2d). MSI status and mutation type provide clinically actionable patient stratification biomarkers. Quantitative therapeutic window analysis combined with PRISM drug screening and structural druggability assessment nominates ARID1A→ARID1B as a leading computational candidate for preclinical development, pending experimental validation.

This work provides a reproducible computational framework for prioritizing context-specific paralog dependencies across solid tumor models and nominates ARID1A→ARID1B as a leading candidate for experimental validation. We provide a complete open-source toolkit, the `paralogSL` R package and a fully reproducible Python pipeline, that requires only publicly available DepMap data and no specialized computational infrastructure.

Methods

Data sources and preprocessing

DepMap 26Q1. CRISPRGeneEffect.csv (420 MB; Chronos-normalized gene dependency scores (hereafter “gene effect scores”) for 1,208 cell lines $\times$ 18,531 genes), OmicsSomaticMutations.csv (554 MB), OmicsExpressionProteinCodingGenesTPMLogp1.csv (483 MB), and Model.csv (681 KB) were downloaded from the DepMap Portal (<https://depmap.org/portal/download/>, release 26Q1) [20, 21]. Damaging mutations were defined as: Nonsense_Mutation, Frame_Shift_Del, Frame_Shift_Ins, Splice_Site, and Translation_Start_Site [22]. Cell lines without expression or dependency data were excluded, yielding 1,208 cell lines. Cross-study dependency data integration followed established methods [31].

HGNC gene families. The HGNC complete set was downloaded from <https://storage.googleapis.com/public-download-files/hgnc/tsv/tsv/>. Protein-coding genes with “Approved” status and non-empty gene_group_id were retained (11,050 genes). All pairwise gene group combinations yielded 66,595 paralog pairs.

CPTAC proteomics. Ovarian HGSC global proteome data were obtained from the NCI Proteomic Data Commons (PDC study PDC000360, $n = 232$) [30]. Seven additional CPTAC cohorts were retrieved via the cBioPortal API [23, 24]: BRCA ($n = 122$), COAD ($n = 95$), LUAD ($n = 84$), GBM ($n = 99$), PDAC ($n = 81$), UCEC ($n = 83$), and LUSC ($n = 108$). Pearson correlations were computed on $\log_2$ -transformed abundance values (≥ 10 paired samples per test).

PRISM drug sensitivity. The PRISM Repurposing secondary screen dataset (1,482 compounds $\times$ 727 cell lines) was downloaded from the DepMap Portal [25]. $\Delta\text{AUC} = \text{mean}(\log_2\text{AUC}_{\text{MUT}}) - \text{mean}(\log_2\text{AUC}_{\text{WT}})$. Selective drugs: BH $q < 0.25$, $\Delta\text{AUC} < 0$, enrichment > 0.1 .

TCGA survival data. Pan-cancer overall survival and gene expression data were obtained from the UCSC Xena browser [26]. Breast cancer patients ($n = 1,082$) were stratified by median paralog expression. Cox proportional hazards models estimated hazard ratios.

Delta Dependency computation

DD was computed as defined in Equation 1. Welch’s t -test (not assuming equal variances). Cohen’s d : standardized mean difference using pooled SD. Minimum 3 mutant and 3 wild-type cell lines per driver. Benjamini-Hochberg FDR correction [29] ($q < 0.20$, permissive stage-1 threshold) was applied *within each driver gene*: for each driver D , all (D, P_i, c_j) combinations (paralog $\times$ cancer type) were corrected jointly, controlling for the total number of paralog pairs tested for that driver. We note that this constitutes a discovery-stage analysis: global FDR correction across all 66,595 paralog pairs $\times$ 40 drivers $\times$ 23 cancer types was not performed, as the goal was to maximize sensitivity for candidate identification rather than to control family-wise error rate. All reported candidates should therefore be considered hypothesis-generating until experimentally validated. Sensitivity analysis showed that at a stricter threshold of $q < 0.10$, PIK3CA $\rightarrow$ PIK3CB remained significant ($q = 0.086$ in Endometrial), while ARID1A $\rightarrow$ ARID1B, despite having the largest $|\text{DD}|$ (0.30), did not survive within-driver FDR correction ($q = 0.50$), reflecting the large number of competing paralog pairs for ARID1A rather than weak signal strength.

Benchmark comparison

Gold-standard positives and matched negatives. Twelve paralog-SL pairs were curated from published experimental studies (genetic screens, viability assays, biochemical validation) independently of any DepMap-derived results (Supplementary Table S3). Of these, 10 are true sequence paralogs (shared gene family, conserved domains), 1 is a partial homolog (STK11→SIK1), and 1 is a functional analog (BRCA1↔BRCA2; included as a mechanistic comparator). Negative pairs comprised all remaining HGNC-defined paralog pairs involving the same 40 driver genes that did not appear in the positive set. This yields an imbalanced evaluation set (12 positives vs. ~194 negatives per lineage), which is representative of the real-world class imbalance but may inflate AUROC; we therefore also report AUPRC where applicable. Directional pairs (e.g., BRCA1→BRCA2 and BRCA2→BRCA1) are scored separately when both driver genes are in the 40-gene panel.

Contextual comparison with published methods.

DD was compared against eight published SL prediction methods evaluated in the CV3 setting by Feng et al. [12]. This comparison is restricted to the paralog-SL prediction task using the 12 gold-standard paralog-SL pairs as the positive set; published method AUROCs are from the same CV3 evaluation framework. AUROC was computed using scikit-learn 1.3.0 [33]. Bootstrap 95% CI: 1,000 iterations, resampling paralog pairs with replacement. For the DD + ID_{≥30%} subset (10 pairs), bootstrap CI was computed separately to account for the reduced sample size. Negative controls: 100 permutations of shuffled known-SL labels.

Sequence identity computation

K-mer-based Jaccard similarity ($k = 3$) was used as a computationally efficient proxy for global protein sequence identity. This metric was validated against Needleman-Wunsch global alignment identity in a subset of 50 protein pairs (13 known paralog pairs and 37 non-paralog controls; Pearson $r = 0.88$, $p < 10^{-16}$; Supplementary Fig. S9). Within-group correlations were $r = 0.91$ (paralogs, $n = 13$) and $r = 0.41$ (non-paralogs, $n = 37$), confirming that the high overall r is not solely driven by between-group variance. The $\geq 30\%$ threshold should therefore be interpreted as an approximate filter for enriching high-identity pairs rather than a precise sequence identity cutoff.

Potential confounders and causal interpretation

DD quantifies an observational association between driver mutation status and paralog dependency, not a causal effect. Potential confounders include: (1) lineage differences between mutant and wild-type cell lines within a cancer type; (2) co-occurring mutations that may independently affect paralog dependency; (3) copy number alterations at paralog loci. We mitigated these through: within-lineage analysis (DD computed separately per cancer type), Welch's t -test (robust to unequal variances), and CNV independence verification (Supplementary Fig. S4, all $R^2 < 0.10$). Benjamini-Hochberg correction was applied within each driver gene to control false discovery rate. Nevertheless, experimental validation (e.g., CRISPR double-knockout in isogenic cell lines) remains essential to establish causality.

CPTAC proteomic co-variation

Protein abundance data for seven CPTAC cohorts (BRCA, COAD, LUAD, GBM, PDAC, UCEC, LUSC; $n = 672$ total) were retrieved via the cBioPortal API [23, 24]. For each paralog pair, Pearson correlation of protein $\log_2$ -abundance was computed across all samples within each cohort. Significance was assessed using BH correction within each pair across 7 cohorts. Mutation-conditioned proteomic analysis was performed for the ARID1A→ARID1B pair in the UCEC cohort: ARID1A mutation status was retrieved from the cBioPortal mutations API (ucec_cptac_2020_mutations), and ARID1B protein abundance was compared between ARID1A-mutant ($n = 37$) and wild-type ($n = 58$) samples using the Wilcoxon rank-sum test. Extension to additional pairs and cohorts is a priority future direction.

Pan-cancer lineage definitions

Twenty-seven solid tumor types defined by matching OncotreePrimaryDisease annotations to curated keyword lists (≥ 6 cell lines with available data). Forty pan-cancer driver genes were selected from TCGA pan-cancer analyses [27] and the COSMIC Cancer Gene Census [28], filtered to genes with ≥ 3 mutant cell lines in at least one lineage (Supplementary Table S5).

MSI stratification

MSI-H was approximated by the presence of damaging mutations in MMR genes (MLH1, MSH2, MSH6, PMS2) or POLE as a proxy for microsatellite instability status; this mutation-based proxy may not perfectly correspond to clinical MSI testing (PCR or IHC). DD computed separately within each subgroup for colorectal and endometrial cancers.

Mutation type classification

Driver mutations classified as truncating (frameshift, nonsense, splice-acceptor, splice-donor, start-loss, stop-loss) or missense (missense_variant, protein_altering_variant, excluding truncating) using VEP VariantInfo annotations. Per-cell-line classification used the most severe mutation type.

Therapeutic window analysis

TI defined as in Equation 2. The denominator uses $\max(|\mu_{\text{CERES}}|, f_{\text{pan-essential}}, 0.01)$ to capture the most conservative estimate of baseline paralog essentiality: $|\mu_{\text{CERES}}|$ reflects mean dependency magnitude (dimensionless, normalized by Chronos), $f_{\text{pan-essential}}$ reflects the fraction of lines where the paralog is essential (dimensionless, range 0–1), and the 0.01 floor prevents division by near-zero values for non-essential paralogs. All three quantities are dimensionless and on comparable scales (0–1); the max operator selects whichever indicates the greatest baseline toxicity risk. Paralogs classified: PAN_ESSENTIAL ($f_{\text{pan-essential}} > 0.5$), HIGH_SELECTIVITY (selectivity > 0.15 and TI > 1.0), MODERATE, or LOW_SELECTIVITY. Analysis across 5 cancer contexts (Ovarian, Endometrial, Breast, Colorectal, PanCancer).

Structural druggability

Sequence-based features (length, MW, pI, GRAVY hydrophobicity, disorder propensity, lysine/cysteine density) from UniProt canonical sequences. Structural similarity: $1 - \overline{|\Delta\text{feature}|}$. Domain architectures from UniProt annotations. PROTAC suitability: $\text{lysine_density} \times 2.0 + \text{disorder_propensity} - \text{cysteine_density} \times 0.5$. Composite targetability: therapeutic index (0.40), selectivity (0.30), structural similarity (0.15), druggability (0.15).

Statistical analyses

AUROC via scikit-learn 1.3.0 [33]. Bootstrap CI: 2.5th and 97.5th percentiles from 1,000 iterations. Negative controls: 10,000 label permutations. Analyses in Python 3.10 with numpy 1.24 [35], scipy 1.11 [34], pandas 2.0, and statsmodels 0.14. The companion R package paralogSL provides DD computation, PCS analysis, benchmark comparison, and CPTAC correlation visualization under the MIT license.

Power analysis and test set stratification

Gold-standard pairs were stratified into two evaluation sets: a *Primary set* of 10 true sequence paralogs (shared gene family and conserved domains) and a *Secondary set* of 2 functional analogs (BRCA1 $\leftrightarrow$ BRCA2, STK11 $\rightarrow$ SIK1) included as mechanistic comparators. Main AUROC results report the Primary set; the Full set (12 pairs) is reported for completeness. With 8–10 positive pairs, the permutation test has limited power: at $\alpha = 0.05$, an estimated ≥ 25 validated positive pairs would be needed to achieve 80% power for detecting AUROC = 0.85 against a null of 0.5. To compensate for this limitation, we complement binary classification metrics with effect-size ranking (—DD—, TI, selectivity) and multivariate regression significance (p -values for individual pairs).

Supplementary information

Supplementary Figures S1–S9 and Supplementary Tables S1–S5 are available as a separate PDF file. Supplementary Fig. S1: Panoramic baseline across 23 solid tumor types (cell line counts and tested pairs, colored by AUROC tier). Supplementary Fig. S2: Cross-cancer DD AUROC detailed values. Supplementary Fig. S3: CPTAC per-cohort scatter plots. Supplementary Fig. S4: CNV independence analysis. Supplementary Fig. S5: PRISM drug selectivity full results. Supplementary Fig. S6: Full mutation-type stratification. Supplementary Fig. S7: Therapeutic window full classification. Supplementary Fig. S8: Full bootstrap distribution and negative control. Supplementary Fig. S9: k -mer Jaccard vs. Needleman-Wunsch alignment identity validation ($n = 50$ pairs, Pearson r shown). Supplementary Table S1: Cell line counts per cancer type. Supplementary Table S2: Complete paralog-SL pair analysis results (118 driver $\times$ paralog $\times$ cancer-type associations passing the minimum sample size filter of ≥ 3 mutant and ≥ 3 wild-type cell lines). Supplementary Table S3: Gold-standard paralog-SL pair evidence. Supplementary Table S4: Pan-cancer summary statistics. Supplementary Table S5: Complete 40-gene driver panel with cancer types, selection source, and functional category. Supplementary Table S6: All 24 paralog pairs ranked by preclinical prioritization score with component values.

Declarations

Ethics approval and consent to participate

Not applicable. This study used exclusively publicly available, de-identified data.

Data availability

All data are publicly available: DepMap 26Q1 from <https://depmap.org/portal/download/>; CP-TAC proteomics from <https://cbioportal.org> and <https://pdc.cancer.gov>; TCGA PanCan from <https://xenabrowser.net>; PRISM from <https://depmap.org/portal/prism/>. HGNC gene families from <https://www.genenames.org/>. Processed results are deposited at the GitHub repository below.

Code availability

All analysis code is available under the MIT license at <https://github.com/tjogzt/paralogSL>. The paralogSL R package (v1.0.0) provides a complete workflow for paralog-SL candidate prioritization:

Workflow: DepMap gene-effect matrix $\rightarrow$ `build_mutation_matrix()` $\rightarrow$ `compute_dd()` per driver $\times$ paralog $\times$ lineage $\rightarrow$ `compute_auroc()` against gold standard $\rightarrow$ `compute_therapeutic_window()` for translational ranking $\rightarrow$ visualization via `plot_pancancer_summary()` and `plot_therapeutic_window()`.

Quick start (3 lines):

```
install.packages("paralogSL") # or devtools::install_github("taozhu/paralogSL")
library(paralogSL)
result <- run_paralog_analysis(dep_matrix, mut_matrix, driver="ARID1A", paralog="ARID1B")
```

The package includes built-in datasets (`known_sl_pairs`, `gyn_drivers`, `solid_tumor_summary`), a vignette with worked examples, and 21 documented functions. A frozen release with all processed data tables is archived at Zenodo (DOI to be provided upon acceptance). The complete Python analysis pipeline (`main.py`, `pcs.py`, `pancancer.py`) is included in the same repository for full reproducibility.

Competing interests

The author declares no competing interests.

Funding

This research received no specific grant from any funding agency.

Authors' contributions

T.Z. conceived the study, developed the methodology, performed all computational analyses, wrote the manuscript, and developed the R package.

Acknowledgements

We thank the DepMap, CPTAC, TCGA, PRISM, and cBioPortal teams for making their data publicly available. We also acknowledge the SynLethDB, HGNC, and UniProt teams for curated resources.

References

- [1] O’Neil NJ, Bailey ML, Hieter P. Synthetic lethality and cancer. *Nat Rev Genet.* 2017;18(10):613–623.
- [2] Huang A, Garraway LA, Ashworth A, Weber B. Synthetic lethality as an engine for cancer drug target discovery. *Nat Rev Drug Discov.* 2020;19(1):23–38.
- [3] Lord CJ, Ashworth A. PARP inhibitors: synthetic lethality in the clinic. *Science.* 2017;355(6330):1152–1158.
- [4] Bryant HE, Schultz N, Thomas HD, Parker KM, Flower D, Lopez E, et al. Specific killing of BRCA2-deficient tumours with inhibitors of poly(ADP-ribose) polymerase. *Nature.* 2005;434(7035):913–917.
- [5] Farmer H, McCabe N, Lord CJ, Tutt ANJ, Johnson DA, Richardson TB, et al. Targeting the DNA repair defect in BRCA mutant cells as a therapeutic strategy. *Nature.* 2005;434(7035):917–921.
- [6] Zheng J, Guo N, Papachristodoulou A, Huss M, Mardinoglu A, Huang T, et al. SynLethDB 2.0: a web-based knowledge graph database on synthetic lethality for novel anticancer drug discovery. *Database.* 2022;2022:baac030.
- [7] Chen L, Zhang Y, Wang X, Li J, Liu H, Wang S. SLGNN: synthetic lethality prediction via graph neural network. *Bioinformatics.* 2023;39(11):btad681.
- [8] Das S, Dey A, Mallik S, Zhao Z. DDSL: deep double-strand learning for synthetic lethality prediction. *Bioinformatics.* 2024;40(3):btae105.
- [9] Liu H, Wang X, Zhang Y, Li J, Chen L, Wang S. MVGCN: multi-view graph convolutional network for synthetic lethality prediction. *Brief Bioinform.* 2024;25(1):bbad449.
- [10] Li J, Zhang Y, Wang X, Liu H, Chen L, Wang S. NSF4SL: negative sampling framework for synthetic lethality prediction. *Brief Bioinform.* 2023;24(5):bbad286.
- [11] Zhang Y, Wang X, Li J, Liu H, Chen L, Wang S. SLAST: sequence-learning approach for synthetic lethality prediction. *Genome Biol.* 2024;25:102.
- [12] Feng Y, Long Y, Wang H, Ouyang Y, Li Y, Wu M, et al. Benchmarking machine learning methods for synthetic lethality prediction in cancer. *Nat Commun.* 2024;15:9058.
- [13] Koonin EV. Orthologs, paralogs, and evolutionary genomics. *Annu Rev Genet.* 2005;39:309–338.
- [14] D’Antonio M, Guerra RF, Cereda M, Marchesi S, Montani F, Nicassio F, et al. Systematic analysis of paralog dependency in cancer. *Cell Syst.* 2023;14(8):675–691.

- [15] Hoffman GR, Rahal R, Buxton F, Xiang K, McAllister G, Frias E, et al. Functional epigenetics approach identifies BRM/SMARCA2 as a critical synthetic lethal target in BRG1-deficient cancers. *Proc Natl Acad Sci USA*. 2014;111(8):3128–3133.
- [16] Helming KC, Wang X, Wilson BG, Vazquez F, Haswell JR, Manchester HE, et al. ARID1B is a specific vulnerability in ARID1A-mutant cancers. *Nat Med*. 2014;20(3):251–254.
- [17] Jia S, Liu Z, Zhang S, Liu P, Zhang L, Lee SH, et al. Essential roles of PI(3)K-p110 β in cell growth, metabolism and tumorigenesis. *Nature*. 2008;454(7205):776–779.
- [18] Wee S, Wiederschain D, Maira SM, Loo A, Miller C, deBeaumont R, et al. PTEN-deficient cancers depend on PIK3CB. *Proc Natl Acad Sci USA*. 2008;105(35):13057–13062.
- [19] Taylor SE, O’Connor CM, Wang Z, Shen G, Song H, Leonard D, et al. PPP2R1A mutations cause synthetic lethality with RNR inhibition through a p53-dependent mechanism. *Cancer Res*. 2022;82(4):721–733.
- [20] Cancer Dependency Map Portal. DepMap 26Q1 Public Release. <https://depmap.org>. 2026.
- [21] Meyers RM, Bryan JG, McFarland JM, Weir BA, Sizemore AE, Xu H, et al. Computational correction of copy number effect improves specificity of CRISPR-Cas9 essentiality screens in cancer cells. *Nat Genet*. 2017;49(12):1779–1784.
- [22] Dempster JM, Boyle I, Vazquez F, Root DE, Boehm JS, Hahn WC, et al. Chronos: a cell population dynamics model of CRISPR experiments that improves inference of gene fitness effects. *Genome Biol*. 2021;22:343.
- [23] Cerami E, Gao J, Dogrusoz U, Gross BE, Sumer SO, Aksoy BA, et al. The cBio cancer genomics portal: an open platform for exploring multidimensional cancer genomics data. *Cancer Discov*. 2012;2(5):401–404.
- [24] Gao J, Aksoy BA, Dogrusoz U, Dresdner G, Gross B, Sumer SO, et al. Integrative analysis of complex cancer genomics and clinical profiles using the cBioPortal. *Sci Signal*. 2013;6(269):p11.
- [25] Corsello SM, Nagari RT, Spangler RD, Rossen J, Kocak M, Bryan JG, et al. Discovering the anticancer potential of non-oncology drugs by systematic viability profiling. *Nat Cancer*. 2020;1(2):235–248.
- [26] Liu J, Lichtenberg T, Hoadley KA, Poisson LM, Lazar AJ, Cherniack AD, et al. An integrated TCGA pan-cancer clinical data resource to drive high-quality survival outcome analytics. *Cell*. 2018;173(2):400–416.e11.
- [27] Bailey MH, Tokheim C, Porta-Pardo E, Sengupta S, Bertrand D, Weerasinghe A, et al. Comprehensive characterization of cancer driver genes and mutations. *Cell*. 2018;173(2):371–385.e18.
- [28] ICGC/TCGA Pan-Cancer Analysis of Whole Genomes Consortium. Pan-cancer analysis of whole genomes. *Nature*. 2020;578(7793):82–93.
- [29] Benjamini Y, Hochberg Y. Controlling the false discovery rate: a practical and powerful approach to multiple testing. *J R Stat Soc B*. 1995;57(1):289–300.

- 612 [30] Edwards NJ, Oberti M, Thangudu RR, Cai S, McGarvey PB, Jacob S, et al. The CPTAC data portal:
613 a resource for cancer proteomics research. *J Proteome Res.* 2015;14(6):2707–2713.
- 614 [31] Pacini C, Dempster JM, Boyle I, Gonçalves E, Najgebauer H, Karakoc E, et al. Integrated cross-
615 study datasets of genetic dependencies in cancer. *Nat Commun.* 2021;12:1661.
- 616 [32] Bitler BG, Aird KM, Garipov A, Li H, Amantangelo M, Kossenkov AV, et al. Synthetic lethality by
617 targeting EZH2 methyltransferase activity in ARID1A-mutated cancers. *Nat Med.* 2015;21(3):231–
618 238.
- 619 [33] Pedregosa F, Varoquaux G, Gramfort A, Michel V, Thirion B, Grisel O, et al. Scikit-learn: machine
620 learning in Python. *J Mach Learn Res.* 2011;12:2825–2830.
- 621 [34] Virtanen P, Gommers R, Oliphant TE, Haberland M, Reddy T, Cournapeau D, et al. SciPy 1.0:
622 fundamental algorithms for scientific computing in Python. *Nat Methods.* 2020;17(3):261–272.
- 623 [35] Harris CR, Millman KJ, van der Walt SJ, Gommers R, Virtanen P, Cournapeau D, et al. Array
624 programming with NumPy. *Nature.* 2020;585(7825):357–362.