

Delta Dependency Profiles Paralog Compensation

Signals Across Solid Tumor Types

Qingqing Mo^{1,2}, Pingbo Chen^{1,2}, Cheng Xu^{1,2}, Ya Wang^{1,2}, Ting Hu^{1,2}, Qian Sun^{1,2}, Tao
Zhu^{1,2,*}

¹ Department of Obstetrics and Gynecology, National Clinical Research Center for Obstetrics and Gynecology, Tongji Hospital, Tongji Medical College, Huazhong University of Science and Technology, 1095 Jiefang Avenue, Wuhan 430030, China

² Key Laboratory of Cancer Invasion and Metastasis (Ministry of Education), Hubei Key Laboratory of Tumor Invasion and Metastasis, Tongji Hospital, Tongji Medical College, Huazhong University of Science and Technology, 1095 Jiefang Avenue, Wuhan 430030, China

* Correspondence: zhutao@tjh.tjmu.edu.cn

Email addresses: qingqingmo520@tjh.tjmu.edu.cn (Q. Mo); supercpb520@163.com (P. Chen); watt15629030676@tjh.tjmu.edu.cn (C. Xu); misswangya@hotmail.com (Y. Wang); huting.tj@163.com (T. Hu); sunqian@tjh.tjmu.edu.cn (Q. Sun); zhutao@tjh.tjmu.edu.cn (T. Zhu)

Abstract

Background: Paralog compensation is a specific synthetic-lethal (SL) mechanism: when a tumor-suppressor driver is lost to mutation, its sequence-similar paralog can become conditionally essential. Systematic prioritization of paralog-SL candidates remains limited; existing models integrate multiple features, whereas a measurable dependency-shift statistic offers a complementary, transparent strategy.

Results: We introduce Delta Dependency (DD), an interpretable metric measuring the dependency shift on a paralog between driver-mutant and wild-type cell lines in DepMap (1,208 lines, 23 solid tumor types). On twelve curated, evidence-tiered pairs in the primary literature-derived benchmark (three gynecological lineages), signed DD achieved AUROC = 0.629 in a positive-unlabeled, internal-consistency evaluation that does not reach conventional statistical significance (pair-clustered bootstrap 95% CI 0.253–0.933; permutation $p = 0.122$; power 0.34 at the observed effect size; 0.613 excluding two DepMap-labelled pairs; 0.525 for the eight pairs with pre-DepMap experimental evidence; the sole lineage-evaluable Tier A/B pair, ARID1A→ARID1B, ranked above all unlabeled controls in both evaluable lineages). As a fully external hold-out, mutation-agnostic |DD|—which does not test the directional, mutation-conditioned proposition that defines signed DD—distinguished screen hits from non-hits in two independent combinatorial CRISPR screens (Harle et al. 2025: AUROC = 0.619, permutation $p = 2 \times 10^{-4}$; Flister et al. 2025: AUROC = 0.696,

$p = 1 \times 10^{-4}$), though neither dataset contained sufficient driver-mutant models to adjudicate the directional hypothesis (0/9 and 2/6 mutant-model observations respectively). Exploratory analysis of seven CPTAC cohorts ($n = 771$) detected protein-level paralog co-variation. Dependency-window scoring recovered the established ARID1A→ARID1B pair (DWS = 2.82) as the highest-selectivity established candidate and nominated KMT2D→KMT2C as the highest-ranked non-benchmark candidate for experimental follow-up.

Conclusions: DD traces each candidate nomination to a measured dependency shift rather than a black-box prediction. The primary internal benchmark does not reach statistical significance (AU-ROC = 0.629, permutation $p = 0.122$), and the paralog’s baseline essentiality (necessity alone, AU-ROC = 0.642) matches DD without requiring mutation-conditioning, underscoring that the current evidence for mutation-specific prioritization is descriptive rather than conclusive. As a fully external test, mutation-agnostic dependency displacement recapitulated digenic signal in two independent screens (Harle 2025, Flister 2025). DD is an association-based prioritization metric, not a causal test of synthetic lethality; all candidates require experimental validation. An open-source R package (paraLogSL) and reproducible pipeline are provided.

Keywords: Synthetic lethality, paralog compensation, DepMap, CPTAC proteomics, Delta Dependency, dependency window

Background

The PARP inhibitor paradigm illustrates the appeal of synthetic lethality (SL): BRCA1/2-mutant tumors depend on PARP-mediated repair, so PARP inhibition kills tumor cells while sparing normal cells [1–5]. Extending this principle to other mutations has proven difficult [6]. Available SL catalogues contain tens of thousands of heterogeneous, unevenly validated interactions [7, 8], and the existing computational prediction methods (graph neural networks [9, 10], graph convolutional networks with dual dropout [11], graph-regularized matrix factorization [12], knowledge-graph neural networks [13], and negative-sample-free contrastive learning [14]) train almost entirely on these limited known pairs. Feng et al. [15] showed that these methods degrade sharply from the CV1 setting (AUROC > 0.90 for the best methods) to the CV3 gene-pair isolation setting, where most fall to near-random performance (AUROC $\approx$ 0.5) and the best published result reaches 0.790 [16], precisely the setting where predictive models matter: nominating genuinely new SL pairs.

Paralog compensation provides a simpler route into this problem [17, 18]. Large-scale analyses of CRISPR dependency data have since established paralog pairs as the most enriched class of synthetic-lethal interactions and shown that interpretable models built on protein-interaction and evolutionary-

conservation features can predict them [19]; combinatorial knockout screens have begun to map paralog synthetic lethality directly [20–25]. Gene duplication events leave behind paralogs: sequence-similar gene pairs with overlapping biochemical roles. When a paralog acting as a tumor suppressor (the “driver”) is inactivated by mutation, its partner paralog can become conditionally essential. This is *sequence-paralog SL*: synthetic lethality between two genes that share high protein sequence identity (operationalized with a validated *k*-mer identity-enrichment proxy; Methods) and conserved functional domains. Known examples include *SMARCA4*→*SMARCA2* (SWI/SNF ATPase subunits) [26], *ARID1A*→*ARID1B* (BAF complex subunits) [27], *AKT1*→*AKT2* [28], and *CDK4*→*CDK6* and *MAP2K1*→*MAP2K2* [20]. Other catalogued pairs carry indirect or direction-restricted evidence: for *EP300* and *CREBBP* (histone acetyltransferases), direct experiments established the reciprocal *CREBBP*→*EP300* direction [29, 30]; for *PIK3CA*→*PIK3CB* (PI3K catalytic isoforms), the primary evidence supports *PTEN*→*PIK3CB* [31]; and *CCNE1*→*CCNE2* is supported by developmental redundancy in mouse double knockouts [32]. Because the evidence behind catalogued pairs varies in both type and direction, we verified each primary citation and assigned every pair an evidence tier (Supplementary Table S3; Results; *CCNE1*→*CCNE2*, for example, is not evaluable under our mutation rules because *CCNE1* is amplification-driven and carries no qualifying hotspot mutations). In every case, the driver and paralog share sequence homology and biochemical function; when the driver is inactivated by mutation, the remaining paralog becomes the cancer cell’s dependency.

Two related phenomena need to be separated from sequence-paralog SL. *Functional analogs* such as *BRCA1* and *BRCA2* share pathway roles (homologous recombination repair) but have no sequence homology. *Pathway-level SL* interactions, such as *BRCA1*↔*PARP1* and *PTEN*↔*PIK3CB* [31, 33], exploit indirect pathway wiring rather than direct biochemical redundancy. Our analysis concentrates on sequence-paralog pairs; functional analogs including *BRCA1/BRCA2* are retained as mechanistic comparators (specificity references, not verified negatives) and flagged throughout (Supplementary Table S3).

D’Antonio et al. published the first systematic analysis of negative genetic interactions between functional paralogs in cancer cell lines, identifying buffering relationships at the gene-family level [18]. That study opened the field but did not benchmark paralog-based predictions against published SL methods, did not incorporate proteomic validation, and did not address clinical stratification or therapeutic window assessment. Complex multi-feature models resist direct experimental traceability: when a neural network predicts SL, reconstructing why a particular pair was ranked highly is difficult, which complicates validation design. A related challenge is that moving from a ranked list to a prioritized candidate set requires conservative integration of orthogonal evidence, yet the optimal readout for paralog compen-

sation is itself unsettled—RNA-level and protein-level signals can diverge [34], and how protein-level compensation relates to dependency-based prioritization remains unexplored.

We defined Delta Dependency (DD) as the difference in mean Chronos gene-effect scores between driver-mutant and wild-type cell lines within a cancer lineage. A positive DD indicates stronger para-log dependency in mutant cells—for ARID1A→ARID1B in ovarian cancer, $|DD| = 0.386$, meaning the paralog’s essentiality score shifts by 0.386 units when ARID1A is mutant. We computed DD across 23 solid tumor types using a 40-gene driver panel (Supplementary Table S4) curated from TCGA PanCanAtlas [35] and the COSMIC Cancer Gene Census [36]. Genomic predictions were supported by CPTAC proteomics (seven cohorts) and PRISM drug screening (1,482 compounds), with MSI-aware cell-line stratification (MSIsensor2 annotation), mutation-type analysis, dependency-window quantification, and exploratory structure-derived targetability descriptors. The primary analysis assessed whether signed DD enriches literature-supported mutation-conditioned paralog dependencies relative to unlabeled para-log pairs.

Results

DD defines a dependency-shift metric for paralog-SL prioritization

We defined Delta Dependency as:

$$DD(D, P, c) = \bar{G}_{P,D-WT,c} - \bar{G}_{P,D-MUT,c} \quad (1)$$

where $\bar{G}$ denotes the mean Chronos gene-effect score, D is the driver gene, P the candidate paralog, and c the cancer lineage. Because lower Chronos gene-effect scores indicate stronger dependency, positive DD indicates increased dependency on paralog P in driver-mutant cell lines relative to wild-type within lineage c , consistent with paralog compensation. Throughout, DD is used as a signed quantity: only positive shifts (increased dependency in driver-mutant cells) are treated as compensation signal, and reverse-direction shifts score against it; analyses scored by $|DD|$, which also credit reverse-direction shifts, are reported as sensitivity analyses.

DD was computed for 66,595 HGNC paralog pairs across 1,208 DepMap cell lines; benchmark evaluation on the curated evaluation set (Results, benchmark comparison) showed signed DD exceeding AUROC = 0.7 in 1 of 8 evaluable lineages on the primary frame (Fig. 1a; Supplementary Table S1): Colorectal (AUROC = 0.775, $n = 62$ pairs) led, followed by Esophagogastric (0.674, $n = 50$ pairs),

non-small cell lung cancer (NSCLC; 0.664), small cell lung cancer (SCLC; 0.651), Ovarian (0.611), Bladder Urothelial (0.531), Endometrial (0.519), and Breast (0.150; per-lineage cell-line counts, tested pairs, and DD AUROC values in Supplementary Fig. S1). Under the relaxed ≥ 3 -per-group sensitivity frame (12 evaluable lineages), 6 of 12 exceeded 0.7, with Biliary Tract (0.990), Pancreatic (0.962), and Melanoma (0.840) the strongest performers. Scoring by $|DD|$, which also credits reverse-direction shifts, yields higher values (7 of 8 primary-frame lineages above 0.7) but is reserved for sensitivity analysis because reverse-direction shifts are not paralog compensation (Methods). Performance did not differ significantly by oncogenic mechanism (Fig. 1b; sensitivity frame, because the primary frame retains only one oncogene-driven lineage): TSG-driven cancers had a mean AUROC of 0.649 ($n = 9$ lineages) versus 0.834 for oncogene-driven cancers ($n = 3$ lineages; permutation $p = 0.127$, exact Mann-Whitney $p = 0.145$). With only three oncogene-driven lineages the comparison is severely underpowered, and we report it as descriptive.

We compared DD against eight published SL prediction methods in the CV3 setting reported by Feng et al. (2024) [15] (Supplementary Table S12). DD achieved AUROC = 0.629 without any training, below the best published result in that setting (SLMGAE [16], 0.790), though we stress that the evaluation frameworks are analogous, not identical. We therefore treat the published values as contextual reference points rather than a formal benchmark. Restricting DD to the three pairs with $\geq 30\%$ sequence identity (two positives) ranked both known positives above the single unlabeled control; with only three pairs this is uninformative and is reported in Supplementary Table S12 without a significance claim. For a true head-to-head comparison on the same 72 paralog pairs (6 known positives), we trained four standard classifiers, namely logistic regression (LR), random forest (RF), SVM-RBF, and SVM-Linear, on five features ($|DD|$, the paralog compensation score ($PCS = \max(\Delta\text{Expression}, 0) \times \max(\text{necessity}, 0)$, where $\text{necessity} = -\overline{\text{Chronos}}$ so that larger values mark more essential paralogs), $|\Delta\text{Expression}|$, necessity, and mutation frequency; the composite score was excluded from classifier inputs because it is a deterministic function of the other features) under leave-one-pair-out cross-validation. The composite score (AUROC = 0.831) and SVM-RBF (0.841) performed best, with no detectable difference between them at this sample size (paired-bootstrap $\Delta\text{AUROC} = -0.01$, 95% CI -0.21 to $+0.18$; the interval is too wide to establish equivalence), followed by random forest (0.722) and DD alone (0.672 on this mean-aggregated pair frame; see “Evaluation frameworks”); the linear classifiers were unstable (SVM-Linear 0.240, LR 0.136). With only 6 positives, multi-feature training amplifies noise and no classifier result should be over-interpreted. On the same 72-pair frame, the composite score reached AUPRC = 0.356 ($4.3\times$ the baseline prevalence of 0.083), confirming enrichment beyond random expectation. The composite is a

deterministic weighted combination of min-max-normalized PCS (0.50), $|DD|$ (0.20), $-\log_{10}(q)$ (0.15), and mutation frequency (0.15), with weights fixed in the analysis code (`pcs.py`) and not fitted by nested resampling; because the same benchmark informed its development, its AUROC is descriptive context rather than independent validation. Min-max scaling of the composite components is computed within each cross-validation fold on training pairs only; this fold-internal scaling leaves the composite AUROC unchanged to three decimals (0.831) and changes AUPRC from 0.356 to 0.368. The composite and the dependency-window score use magnitude-based $|DD|$ as an effect-size term, whereas signed DD remains the directional metric for biological interpretation. These comparisons are confined to the paralog-based SL task; DD is not intended for pathway-level interactions such as $BRCA1 \leftrightarrow PARP1$. A DD + k -mer identity-enrichment filter ($\geq 30\%$ proxy) ranked both positives above the single negative on a 3-pair subset, but with only two positives this is uninformative (Supplementary Table S12). Unlike neural networks, every DD nomination maps to a single measured quantity, the dependency shift, which makes it straightforward to design a validation experiment.

Component decomposition on the same test set: signed DD (AUROC = 0.629) versus PCS (0.825), Δ Expression alone (0.547), and necessity (0.642; Fig. 1d). A paired bootstrap over entries found no significant difference between DD and any other component (PCS minus DD +0.195, 95% CI -0.132 to +0.545). Of the components, PCS carries substantial signal, which motivates the composite score. Because unlabeled controls are not verified negatives, AUROC estimates are optimistic; under random hidden-positive contamination at $\pi = 5\%$, the primary AUROC attenuates to 0.580 (2,000 seeded relabelings; full contamination analysis in Supplementary Table S13).

We evaluated DD on evidence-tiered test sets (Table S3), reclassifying the twelve curated pairs after verifying each primary citation. Three pairs have direct genetic synthetic-lethal evidence from dual-gene perturbation assays (*Tier A*: $AKT1 \rightarrow AKT2$ [28], $CDK4 \rightarrow CDK6$ and $MAP2K1 \rightarrow MAP2K2$ [20]). Two further pairs are demonstrated selective dependencies in driver-mutant cells with functional validation (*Tier B*: $SMARCA4 \rightarrow SMARCA2$ [26], $ARID1A \rightarrow ARID1B$ [27]); the $Tier A \cup Tier B$ set constitutes the primary literature-derived benchmark. Five pairs rest on indirect, reciprocal-direction-only, or DepMap-derived evidence (*Tier C*: $EP300 \rightarrow CREBBP$ [29, 30], $PIK3CA \rightarrow PIK3CB$ [31], $CCNE1 \rightarrow CCNE2$ [32], $FBXW7 \rightarrow FBXW2$, $PPP2R1A \rightarrow PPP2R1B$), and two serve as mechanistic comparators ($BRCA1 \leftrightarrow BRCA2$, $STK11 \rightarrow SIK1$). On the full curated set of twelve pairs, evaluated in the three gynecological lineages carried through the benchmark pipeline (Ovarian, Endometrial, Cervical), signed DD achieved AUROC = 0.629 (110 driver $\times$ paralog $\times$ lineage entries, 8 positive entries). Label-quality sensitivity analyses: excluding the two DepMap-derived pairs gave AUROC = 0.613, restricting to the eight pairs with

pre-DepMap experimental evidence gave 0.525, and a direction-strict analysis relabelling the EP300→CREBBP entries as non-positive left the estimate unchanged at 0.629. The signed-metric estimate thus attenuates when DepMap-era labels are removed, and we report this dependence explicitly (Discussion). On the primary Tier A $\cup$ Tier B frame, both evaluable positive entries (the same pair, ARID1A→ARID1B, in Endometrial and Ovarian cancers) ranked above all 102 unlabeled control entries (AUROC = 1.000). With only one biological pair evaluable across two lineages, this estimate is degenerate and reported for completeness only; it does not constitute evidence of predictive performance. SMARCA4→SMARCA2 and the three Tier A pairs had too few qualifying driver-mutant cell lines for lineage-level evaluation under the ≥ 5 -per-group rule. With at most 2–3 strictly validated positives, AUROC computed on small subsets is not informative, so we rely on the full-set estimate and supplement it with effect-size ranking: ARID1A→ARID1B has the largest lineage-level DD (0.386 in Ovarian cancer; Hedges' $g = 1.39$; Supplementary Table S7) and a top-tier dependency-window and selectivity profile (Supplementary Table S9), providing orthogonal prioritization beyond classification performance.

Leave-one-lineage-out AUROC over the three evaluation lineages was stable (range 0.606–0.692). Bootstrap resampling of pairs (1,000 iterations, 72 pairs, 6 positives) on the max-aggregated pair frame gave a 95% CI of 0.294–0.845 (Supplementary Fig. S2). On the primary lineage-level frame, a pair-clustered bootstrap (10,000 iterations, resampling whole pairs with all their lineage entries) gave a wider 95% CI of 0.253–0.933 around the 0.629 estimate, a driver-block bootstrap (10,000 iterations, resampling driver genes together with all their pairs) gave a 95% CI of 0.272–0.929, and pair-level label permutation (10,000 permutations) gave $p = 0.122$; the effective sample size is ≈ 85 of the 110 entries after accounting for within-pair correlation. In the pair-level aggregated frameworks, null distributions from 10,000 label permutations were centered at 0.501 and neither aggregated estimate significantly exceeded its null (max aggregation 0.598, empirical $p = 0.223$; mean aggregation 0.672, $p = 0.090$; Supplementary Fig. S2). With only 8 positive entries (6 positive pairs) the evaluation is underpowered (at the observed 0.629, power to reject a null of 0.5 is 0.34, and ≥ 45 validated positive pairs would be required for 80% power at $\alpha = 0.05$; see Methods), so all benchmark AUROCs are reported as point estimates with explicit uncertainty, with the lineage-level frame retained as the primary framework throughout. Copy number alterations explained $< 10\%$ of gene-effect variance for all 23 paralog genes tested (maximum $R^2 = 0.057$; Supplementary Fig. S3), and multivariate regression controlling for CNV left DD estimates virtually unchanged (mean $|\Delta DD| = 0.005$); because Chronos scores already model out screen-level copy-number bias, this adjustment targets residual paralog-specific confounding. ARID1A→ARID1B remained highly significant after CNV adjustment (adjusted $p = 6.5 \times 10^{-42}$), as

213 did SMARCA4→SMARCA2 ($p = 3.5 \times 10^{-29}$) and EP300→CREBBP ($p = 1.4 \times 10^{-13}$); in a dataset
 214 of this size such p-values are dominated by sample size and residual confounding rather than effect
 215 magnitude, so these adjustments are interpreted as negative controls on the effect estimates (mean
 216 $|\Delta DD| = 0.005$), not as additional statistical evidence. Controlling for expression ($G_P \sim D_{\text{mutation}} +$
 217 Expr_P) produced similar results (mean $|\Delta DD| = 0.014$); significance was retained by the strongest
 218 pairs (ARID1A→ARID1B $p = 4.1 \times 10^{-41}$, PIK3CA→PIK3CB $p = 1.5 \times 10^{-11}$, EP300→CREBBP
 219 $p = 6.3 \times 10^{-13}$), whereas weaker pairs did not reach significance in the pan-cancer frame, consis-
 220 tent with lineage-specific compensation. TP53 co-mutation, present in 76/121 ARID1A-mutant lines,
 221 had negligible impact on ARID1A→ARID1B ($\Delta DD = 0.0004$, adjusted $p = 5.8 \times 10^{-42}$). Adding lin-
 222 eage fixed effects attenuated but preserved the association (lineage-adjusted $p = 2.2 \times 10^{-13}$; TP53-
 223 and CNV-adjusted with lineage fixed effects $p = 7.0 \times 10^{-13}$). Re-estimating these models with stan-
 224 dard errors clustered by Oncotree lineage likewise preserved significance for all three headline pairs
 225 (ARID1A→ARID1B $p = 1.7 \times 10^{-9}$; SMARCA4→SMARCA2 $p = 5.2 \times 10^{-3}$; EP300→CREBBP
 226 $p = 3.2 \times 10^{-6}$; Supplementary Table S11).

227 **External digenic screens.** As a fully external test we applied the frozen pipeline to two recent
 228 combinatorial CRISPR screens that played no role in method development, feature choice, or candidate
 229 selection [24, 25]: the directed screen of Harle et al. (472 nominated pairs across 27 melanoma, lung, and
 230 pancreatic cell lines; 117 experimentally determined synthetic-lethal pairs) and the unbiased paralogue-
 231 wide screen of Flister et al. (36,648 pairs screened in NCI-H1299 and MDA-MB-231; 146 lethal pairs).
 232 Scoring each pair by mutation-agnostic $\max|\Delta DD|$ (Methods), dependency displacement distinguished
 233 screen hits from non-hits in both datasets (Harle: AUROC = 0.619, permutation $p = 0.0002$, 411 evalu-
 234 able pairs including 98 hits; Flister: AUROC = 0.696, $p = 0.0001$, 29,457 evaluable pairs including 114
 235 hits; per model 0.671 and 0.744; Additional file 1, Supplementary Methods). Dependency displacement
 236 therefore carries reproducible signal for digenic synthetic lethality at scale, on experimental platforms
 237 entirely independent of DepMap; the null in4mer comparison (Limitations) reflects its 10-positive sam-
 238 ple size rather than an absence of signal. Two qualifications apply. First, the Harle library was nominated
 239 from prior RNAi and digenic studies and is therefore partially enriched for previously reported candi-
 240 dates; the Flister screen is free of this enrichment. Second, the driver-conditioned, directional proposition
 241 that defines DD could not be adjudicated in either screen, because only three curated pairs per dataset had
 242 any driver-mutant models: pooled digenic hit rates were 0/9 in driver-mutant versus 27/121 in wild-type
 243 Harle models (Fisher $p = 1.0$) and 2/6 versus 59/235 in Flister models ($p = 0.475$). This reflects the rar-
 244 ity of our driver mutations among screened lines and the low genotype penetrance of digenic interactions

245 reported by both studies. The single directionally consistent case, EP300→CREBBP in Flister (hits in
246 2/3 mutant versus 20/46 wild-type models), is anecdotal.




Fig. 1: Delta Dependency: pan-cancer performance, benchmark, and validation. **a**, DD AUROC across 8 evaluable solid tumor types on the primary ≥ 5 -per-group frame (red: AUROC > 0.7 ; blue: 0.5 – 0.7 ; gray: < 0.5). Dashed line, random classifier (0.5). Bar labels show n tested pairs. Sample sizes and full lineage names in Supplementary Table S1; the ≥ 3 -per-group sensitivity frame yields 12 evaluable lineages. **b**, DD performance stratified by oncogenic mechanism (blue, TSG-driven; red, oncogene-driven). Mean AUROC does not differ significantly between TSG-driven tumors ($n = 9$; mean 0.649) and oncogene-driven tumors ($n = 3$; mean 0.834 ; permutation $p = 0.127$). Points, individual lineages; center line, group mean. Computed on the ≥ 3 -per-group sensitivity frame because the primary ≥ 5 frame retains only one oncogene-driven lineage. **c**, Benchmark selection and evaluation framework. Twelve curated pairs were evidence-tiered after verification of each primary citation (three Tier A, two Tier B, five Tier C, two comparators; Supplementary Table S3); the Tier A $\cup$ Tier B set (five pairs) constitutes the primary literature-derived benchmark. The evaluation frame spans the three gynecological lineages (Ovarian, Endometrial, Cervical) with ≥ 5 mutant and ≥ 5 wild-type cell lines per entry, yielding 110 driver $\times$ paralog $\times$ lineage entries (8 positive; 102 unlabeled controls). Performance is reported under three aggregation frameworks defined before final reporting (see “Evaluation frameworks”). **d**, Component decomposition: signed DD (AUROC = 0.629) versus PCS (0.825), Δ Expression alone (0.547), and necessity (0.642). Inset: paired bootstrap distribution of PCS minus DD (10,000 resamples; solid line, mean $+0.195$; dashed lines, 95% CI -0.132 to $+0.545$; the difference is not significant). The pair-level bootstrap and negative control are in Supplementary Fig. S2.

Proteomic context: paralog protein co-variation

Whether paralog compensation acts at the transcript or protein level matters. We analyzed quantitative proteomics from seven independent CPTAC cohorts: BRCA ($n = 122$), COAD ($n = 97$), LUAD ($n = 110$), GBM ($n = 99$), PDAC ($n = 140$), UCEC ($n = 95$), and LUSC ($n = 108$), totaling 771 tumor samples with matched protein abundance [37–39].

Protein-level paralog co-variation was detectable across cohorts (Fig. 2a). EP300↔CREBBP showed the most consistent signal: significant in 5 of 7 cohorts after within-pair BH correction (BRCA: $r = 0.263$, $q = 0.005$; COAD: $r = 0.510$, $q < 10^{-3}$; LUAD: $r = 0.380$, $q < 10^{-3}$; GBM: $r = 0.501$, $q < 10^{-3}$; LUSC: $r = 0.372$, $q < 10^{-3}$; mean $r = 0.320$; Fig. 2b). Across all 26 pairs $\times$ 7 cohorts, 122 tests were evaluable with ≥ 10 paired protein measurements (of 182 possible), and 42 were significant at global BH $q < 0.05$ (per-cohort scatter plots in Supplementary Fig. S4). KRAS↔HRAS had the strongest single-cohort correlation ($r = 0.667$, $p < 10^{-4}$ in LUSC). Protein co-variation alone does not prove compensatory dependency; it could reflect shared transcriptional regulation, tumor purity, or copy number linkage. These results provide orthogonal contextual evidence of co-variation, not direct validation of conditional dependency.

Matched mRNA co-variation was also detectable. Across 179 pair $\times$ cohort mRNA tests covering the same 26 pairs in the seven cohorts, the median mRNA correlation was $r = 0.074$, versus $r = 0.140$ for the 122 pair $\times$ cohort tests with matched protein measurements; protein correlations were not systematically stronger (sign test on $|r|$, $p = 0.79$), and for EP300↔CREBBP the mRNA correlation exceeded the protein correlation in all seven cohorts (e.g., UCEC $r = 0.765$ vs. $r = 0.081$). Adjusting protein correlations for tumor purity where annotation exists (BRCA, LUAD, LUSC, UCEC; 69 tests) changed them negligibly (median $|\Delta r| = 0.014$; no test lost significance after adjustment), so the protein co-variation is not a purity artifact. As an SL-discrimination feature, however, the RNA layer remained weak: Δ Expression alone achieved AUROC = 0.547 for discriminating known paralog-SL pairs (Fig. 2c), below DD (0.629): a co-variation-versus-discrimination distinction that may contribute to the underperformance of transcript-only predictors. To test directly for compensatory upregulation, we compared ARID1B protein abundance in ARID1A-mutant ($n = 37$) versus wild-type ($n = 58$) endometrial tumors in the UCEC cohort. ARID1B abundance was higher in mutant tumors, but the difference was not statistically significant ($\log_2$ difference 0.063, fold change ≈ 1.05 ; bootstrap 95% CI -0.018 to $+0.141$; rank-biserial $r = 0.213$; Wilcoxon $p = 0.082$, BH $q = 0.27$ across the ten evaluable UCEC mutation-conditioned tests; Fig. 2d). The direction of the effect, however, is consistent with a model in which loss of one complex subunit stabilizes its paralog.




Fig. 2: CPTAC proteomic co-variation across seven cohorts. **a**, Cross-cancer paralog protein co-variation (Pearson's r) across seven independent CPTAC cohorts ($n = 771$ total). Significance: *** $q < 0.001$, ** $q < 0.01$, * $q < 0.05$ (BH-corrected within each pair; small in-cell asterisks). Blue cells with significance markers indicate significant *negative* correlations. Large red asterisks at the right margin mark known SL pairs. **b**, Representative scatter plot: EP300 vs. CREBBP protein abundance in the GBM cohort ($n = 99$, $r = 0.501$, $p < 10^{-4}$). Gray band, 95% CI of linear regression. **c**, Dependency-based (signed DD, AUROC = 0.629) vs. RNA-level (Δ Expression, AUROC = 0.547) prediction performance for known paralog-SL pairs. The RNA signal is weak (dashed line, random expectation). **d**, Mutation-conditioned proteomics: ARID1B protein abundance in ARID1A-mutant ($n = 37$) vs. wild-type ($n = 58$) endometrial tumors (UCEC CPTAC cohort). ARID1B was higher in ARID1A-mutant tumors (fold change ≈ 1.05 ; Wilcoxon $p = 0.082$, not significant), a directionally concordant but statistically inconclusive signal.

DD signal across clinical contexts: MSI status and mutation type

Microsatellite instability (MSI) is a routine clinical biomarker in colorectal and endometrial cancers, and MSI-H tumors accumulate hypermutation-driven passenger events. We tested whether this mutational background dilutes paralog compensation signals by classifying cell lines with the official DepMap mi-

crosatellite annotation (MSIsensor2 MSIscore; MSI-H = score > 20; Methods) and computing DD within each subgroup. On the ≥ 3 -per-group sensitivity frame (Fig. 3a), MSI-H and MSS subgroups showed comparable paralog-SL predictive signal in both cancer types: endometrial AUROC = 0.303 (MSI-H, $n = 17$) vs. 0.500 (MSS, $n = 11$), $\Delta = -0.197$; colorectal AUROC = 0.558 (MSI-H, $n = 14$) vs. 0.545 (MSS, $n = 45$), $\Delta = +0.013$. On the primary ≥ 5 frame the endometrial subgroups are not evaluable (fewer than two curated positive pairs) and the colorectal subgroups rest on two positives each (0.574 MSI-H vs. 0.298 MSS), so the MSI contrast is strictly exploratory. Subgroup sizes are modest and only 3–4 curated positive pairs were evaluable per subgroup; a stratified bootstrap test (10,000 resamples preserving within-subgroup label structure) found no significant MSI-H versus MSS difference in either cancer type (endometrial Δ AUROC = -0.197 , bootstrap $p = 0.534$; colorectal $\Delta = +0.013$, $p = 0.967$). A higher-powered formulation avoids subgrouping altogether: per-pair ordinary least-squares models of paralog gene effect on driver mutation $\times$ MSI-H status, fitted on all annotated lines of each lineage (colorectal $n = 63$, endometrial $n = 28$; Methods), tested directly whether DD itself varies with MSI background. Across 47 evaluable driver $\times$ paralog pairs in colorectal and 24 in endometrial cancer, no pair showed significant MSI modulation after Benjamini–Hochberg correction (minimum $q = 0.157$ and 0.235, respectively); interaction coefficients were small (median $|\Delta DD|$ 0.062 in colorectal and 0.076 in endometrial) and all but four (colorectal) and one (endometrial) of the pair-level 95% confidence intervals spanned zero; the counts of nominally significant interactions did not exceed the 5% chance rate (4 and 1 pairs; binomial $p = 0.207$ and 0.708), and combined tests were null (Stouffer $p = 0.667$ and 0.276). For the lead candidate ARID1A $\rightarrow$ ARID1B, the MSI-H versus MSS difference in DD was negligible in both lineages (endometrial $\Delta DD = +0.02$, $p = 0.96$; colorectal $+0.04$, $p = 0.79$), and a continuous formulation using the MSIsensor2 score gave concordant results (Methods). We therefore find no evidence that a hypermutated background either dilutes or amplifies paralog-SL detection by DD; these analyses remain hypothesis-generating.

Mutation type further modulated the signal. For the leading TSG pair, truncating mutations produced substantially stronger DD than missense variants (ARID1A $\rightarrow$ ARID1B in Ovarian: DD = 0.388 truncating vs. 0.020 missense, $\Delta DD = +0.368$; Fig. 3b), and in Breast cancer truncating-only and missense-only DD performed similarly (AUROC = 0.460 vs. 0.437, both below 0.5; Supplementary Fig. S5). The pattern was not universal, however: across all tested pairs, mean $|DD|$ was modestly larger for missense than for truncating subgroups (0.042 vs. 0.026; paired t -test $p = 3 \times 10^{-4}$), driven mainly by oncogene pairs such as PIK3CA $\rightarrow$ PIK3CB in which hotspot missense mutations define the mutant group (Supplementary Fig. S5b). These are point estimates from small per-subgroup samples and should be treated as

descriptive observations.

As an exploratory clinical correlation, we tested whether paralog expression tracks patient outcome in TCGA. In the BRCA cohort ($n = 1,069$, 151 events), higher *ARID1B* expression was associated with shortened overall survival (multivariable Cox HR = 1.31 per SD, 95% CI: 1.087–1.569, $p = 0.004$, adjusted for age and AJCC stage; Supplementary Fig. S10), though the association reached FDR significance only after covariate adjustment and does not establish a relationship between ARID1B expression and ARID1A-conditioned dependency. Four of the 32 paralog genes passed Benjamini–Hochberg correction (full results in Supplementary Table S8). *SMARCA2* showed no association (HR = 0.96), a reminder that expression-level prognosis and dependency-based synthetic lethality are distinct phenotypes; these associations are hypothesis-generating. Cohort-specific extensions in UCEC and OV are reported in the Supplementary TCGA survival extension. Mutational co-occurrence analysis showed that paralog pairs co-occur rather than being mutually exclusive, but the pan-cancer signal is partly confounded by tumor mutational burden (TMB): after TMB adjustment only ARID1A/ARID1B (adjusted OR = 2.38, $p = 0.002$) and EP300/CREBBP (adjusted OR = 2.14, $p = 0.011$) remained significant (per-lineage analysis in Supplementary Fig. S6). These observations motivate analysing dependency phenotypes directly but do not constitute evidence for synthetic lethality.




Fig. 3: Clinical stratification by MSI status and mutation type. **a**, DD AUROC for MSI-H vs. MSS subgroups in colorectal and endometrial cancers, using the official DepMap MSIsensor2 annotation (MSI score > 20 = MSI-H), computed on the ≥ 3 -per-group sensitivity frame (the primary ≥ 5 frame is not evaluable for endometrial and shows no colorectal difference; see text). Sample sizes shown above bars; hatched bars with *(insuff.) mark subgroups with too few lines for an AUROC estimate. MSI-H and MSS subgroups showed comparable DD signal in both cancer types (stratified bootstrap, not significant; a higher-powered full-cohort interaction regression likewise found no significant MSI modulation of DD after FDR correction; see text). **b**, DD (WT – MUT) for truncating vs. missense mutations in leading TSG paralog-SL pairs across four cancer types. Positive values indicate compensation; larger truncating bars indicate stronger compensation under truncating mutations. See Supplementary Fig. S5 for full results.

331 Dependency-window scoring and pharmacologic context prioritize ARID1A→ARID1B

332 To assess pharmacologic relevance, we analyzed the PRISM Repurposing dataset (1,482 compounds
 333 across 727 cell lines) [40]. At a discovery-stage threshold (BH $q < 0.25$, $\Delta\text{AUC} < 0$, enrichment $>$
 334 0.1; Supplementary Fig. S7), 633 driver-genotype–drug sensitivity associations passed. By definition of
 335 the BH q -value, up to $\sim 25\%$ of associations passing this cutoff are expected to be false positives; the
 336 relaxed threshold is deliberate because this is a screen, not a confirmatory analysis; proliferation rate
 337 was not controlled, so these associations are quality-control context only. Known drug-target biology
 338 was recovered (Supplementary Fig. S7b): MEK inhibitors killed KRAS-mutant lines (AZD8330: $d =$
 339 -0.62 , $\Delta\text{AUC} = -0.37$, $q = 0.0007$; Trametinib: $d = -0.57$, $\Delta\text{AUC} = -0.24$, $q = 0.0006$), mTOR
 340 and AKT inhibitors killed PTEN-mutant lines (Everolimus: $d = -0.61$, $\Delta\text{AUC} = -0.32$, $q = 0.0009$;
 341 Ipatasertib: $d = -0.99$, $\Delta\text{AUC} = -0.14$, $q = 0.0001$), and HDAC inhibitors killed EP300-mutant lines
 342 (Panobinostat in Ovarian cancer: $d = -1.45$, $\Delta\text{AUC} = -0.34$, $q = 0.010$). These associations recover
 343 expected driver-genotype drug biology and serve as an assay-validity control for the PRISM layer; they
 344 do not demonstrate paralog-specific targeting.

To distinguish context-selective paralogs from pan-essential ones, we defined a dependency window score (DWS; an in vitro dependency-selectivity ratio, not a clinical dose-toxicity therapeutic index):

$$\text{DWS} = \frac{\max(\text{DD}, 0)}{\max(|\mu_{\text{Chronos}}|, f_{\text{pan-essential}}, 0.01)} \quad (2)$$

where $f_{\text{pan-essential}}$ is the fraction of cell lines with paralog Chronos score < -0.5 , and selectivity is the mutant-minus-wild-type difference in that fraction between driver-mutant and wild-type groups. The signed numerator credits only compensation-direction shifts; the $|\text{DD}|$ -numerator variant, which also credits reverse-direction shifts, is reported as a sensitivity analysis (Supplementary Table S10). Candidate rankings were stable under alternative denominator constructions (max versus mean, floor values of 0.001–0.05; Supplementary Table S10), indicating that the prioritization is not an artifact of a single DWS parametrization. Each paralog was classified into one of four selectivity tiers (Fig. 4). Two pairs fell into HIGH_SELECTIVITY: ARID1A→ARID1B (mean DWS = 2.82 across four cancer contexts; mean selectivity +0.28) and the established synthetic-lethal pair SMARCA4→SMARCA2 (4.87; +0.18), whose recovery acts as a positive control for the framework. This is what one would expect if ARID1A loss creates a specific dependency on ARID1B-containing BAF complexes [27]: within the dependency-window module ARID1A→ARID1B attained the highest mean selectivity and the highest mean positive DD (0.270) of all evaluated pairs (Supplementary Table S5). NF1→RASA2 attains a moderate signed DWS (1.82) that derives entirely from a single context (Endometrial), where the 0.01 denominator floor is active; its bootstrap rank interval is wide (95% CI 1–18) and selectivity is ≈ 0 , so NF1→RASA2 is not a selective candidate. BRCA1↔BRCA2 and PIK3R1→CRKL were classified as PAN_ESSENTIAL (pan-essential fraction 0.53–0.55 and 0.73), suggesting that targeting them directly would cause broad cellular toxicity. Five paralog pairs maintained DWS > 1.0 in ≥ 2 cancer contexts (Supplementary Fig. S8).

As an exploratory complement, we computed sequence- and structure-derived targetability descriptors for 13 paralog pairs (Supplementary Fig. S9; Methods). Five pairs (38%) showed perfect domain-architecture conservation, and EP300↔CREBBP was the most similar pair (0.976), consistent with shared histone acetyltransferase function. A composite score with heuristically fixed weights (max-normalized DWS 0.40, rescaled selectivity 0.30, targetability 0.30) ranked ARID1A→ARID1B first (0.823), followed by EP300→CREBBP (0.644) and NF1→RASA2 (0.554), making ARID1A→ARID1B the highest-priority established mutation-conditioned pair under the combined criteria (Supplementary Table S9). These descriptors are heuristic proxies, not a formal druggability assessment, and were not

374 calibrated against drug-development data (Methods; Limitations).




Fig. 4: Dependency-window classification. Bubble size, |mean DD|; color, selectivity tier ($n = 21$ paralog pairs). Dashed lines mark the HIGH_SELECTIVITY thresholds (selectivity > 0.15 , DWS > 1.0); NF1→RASA2 has moderate signed DWS (1.82) resting on a floored denominator but near-zero selectivity (0.005), hence MODERATE. See Supplementary Fig. S8 for the full ranking, tier composition, and cross-context consistency, Supplementary Fig. S9 for the exploratory structure-derived descriptors; PRISM pharmacologic context is shown in Supplementary Fig. S7.

375 Discussion

376 We set out to test whether a single, interpretable metric, the dependency shift between driver-mutant
 377 and wild-type cells, could capture the paralog-SL prioritization signal. DD provided a transparent uni-
 378 variate signal, whereas multi-feature performance estimates were unstable because only six positive
 379 pairs were available. On the 72-pair benchmark frame (three gynecological lineages), the interpretable
 380 composite score reached AUROC = 0.831 in head-to-head comparison, with no detectable difference
 381 from the best multi-feature classifier (SVM-RBF, 0.841; classifier range 0.240–0.841; paired-bootstrap
 382 Δ AUROC = -0.01 , 95% CI -0.21 to $+0.18$, too wide to establish equivalence), while DD alone reached
 383 0.672 on the same mean-aggregated pair frame (0.598 with max aggregation; “Evaluation frameworks”).

The dependency shift carries a substantial but not dominant share of the signal; PCS and expression features contribute complementary information. Necessity alone—the paralog’s mean essentiality across all cell lines, without conditioning on driver mutation—reached AUROC = 0.642 on the primary frame, numerically matching signed DD (0.629; paired-bootstrap $\Delta = +0.013$, 95% CI -0.338 to $+0.374$). Adding signed DD to necessity in an equal-weight combination raised AUROC marginally to 0.652 ($\Delta = +0.023$, CI -0.125 to $+0.178$, 61.5% of resamples positive), well within the noise floor of eight positives. DD therefore does not demonstrate a significant incremental contribution beyond baseline essentiality on the primary frame; its value lies instead in its directionality and mechanistic interpretability, which necessity alone cannot provide. The point is most visible where an unconditional essentiality metric is actively misleading: in Breast cancer, where PIK3CA→PIK3CB shows the strongest reverse-direction shift, signed DD yields AUROC 0.150 (below chance), correctly flagging oncogene-addiction displacement rather than compensation, whereas necessity-ranked control pairs would instead reward the very dependencies DD flags as anti-compensation.

Our seven-cohort CPTAC survey found consistent paralog protein co-variation: EP300↔CREBBP was significant in five of seven cohorts, and KRAS↔HRAS reached $r = 0.667$ in LUSC. The protein signal is robust to tumor-purity adjustment (median $|\Delta r| = 0.014$ across 69 tests). Matched mRNA co-variation is also present, however: for EP300↔CREBBP the mRNA correlation exceeded the protein correlation in every cohort, and across the 122 matched pair×cohort tests, protein correlations were not systematically stronger than mRNA correlations (sign test on $|r|$, $p = 0.79$). A pure complex-stoichiometry model, in which loss of one subunit stabilizes its paralog strictly at the protein level, therefore cannot by itself account for the observed co-variation, which more likely reflects co-regulation at both layers; Venkatesh et al. attribute much paralog protein compensation to hemizygous copy-number loss rather than point mutation [34]. What separates the layers in our data is discriminative rather than correlative signal: RNA-level Δ Expression was a weak SL feature (AUROC = 0.547), whereas the dependency-layer statistic (DD) captures post-transcriptional essentiality biology directly. The mutation-conditioned analysis in UCEC was directionally concordant but statistically inconclusive: ARID1B protein abundance was higher in ARID1A-mutant tumors (fold change ≈ 1.05 , $p = 0.082$, BH $q = 0.27$), and the result needs replication in additional cohorts and pairs. Whether protein abundance itself could substitute for dependency data as a prioritization layer is untested, and PROTAC-based strategies that exploit protein abundance rather than transcriptional regulation remain a plausible but untested direction.

Signed-DD performance did not differ significantly between TSG-driven (mean AUROC = 0.649, $n = 9$) and oncogene-driven (0.834, $n = 3$) lineages on the ≥ 3 -per-group sensitivity frame (permu-

tation $p = 0.127$); with only three oncogene lineages the comparison is underpowered, and the primary ≥ 5 frame retains a single oncogene-driven lineage, precluding estimation there. The more informative contrast is at pair level: TSG loss-of-function pairs show the expected compensation direction (positive DD; e.g., ARID1A→ARID1B, DD = 0.386 in Ovarian cancer), whereas the oncogene pair PIK3CA→PIK3CB shows a consistent reverse-direction shift (negative DD in all eight evaluable primary-frame lineages; e.g., DD = -0.258 in Breast cancer). Cells carrying an activated *PIK3CA* allele are less, not more, dependent on the paralogous catalytic subunit, plausibly because the mutant isoform itself sustains downstream signalling. Signed DD makes this mechanistic distinction explicit; scoring by $|\text{DD}|$ would count such anti-dependencies as hits, which is a principal reason we report signed DD as the primary metric.

A full-cohort interaction analysis found no evidence that MSI status modulates DD signal strength in the colorectal or endometrial lineages (Fig. 3a; no pair significant after FDR correction), although power remains limited for rare driver strata, whereas truncating mutations elicited stronger compensation than missense variants for several TSG pairs. These observations generate testable hypotheses for experimental work rather than a clinical enrichment strategy: truncating ARID1A mutations mark the context in which compensation is strongest, and baseline ARID1B abundance could be examined as a correlate of dependency in isogenic models; the data provide no basis for stratifying by MSI status, and this negative result warrants re-estimation in larger cohorts. Our DWS framework nominated ARID1A→ARID1B (DWS = 2.82) as the highest-selectivity established candidate (SMARCA4→SMARCA2 had the higher DWS, 4.87) while flagging BRCA1/2 paralogs as pan-essential. The DWS is based entirely on in vitro data and cannot predict organ-specific in vivo toxicity; natural extensions would incorporate tissue expression data from GTEx, mouse knockout phenotypes from IMPC, and germline constraint metrics (pLI/LOEUF) from gnomAD. The SMARCA4→SMARCA2 precedent illustrates both the feasibility and the demands of such a translational path: selective degradation of SMARCA2 was first achieved with VHL-based PROTAC degraders [41, 42], and the clinical-stage SMARCA2-selective degrader PRT3789 reduced SMARCA2 protein in patients' peripheral blood mononuclear cells and produced RECIST-confirmed partial responses in biomarker-selected SMARCA4-mutated cancers [43]. That path required a dedicated bromodomain ligand, ternary-complex optimization, and extensive medicinal chemistry, none of which our computational targetability descriptors attempt to model [6].

D'Antonio et al. [18] published the foundational analysis of paralog genetic interactions, and De Kegel et al. [19] systematically predicted robust paralog synthetic lethality across hundreds of cancer cell lines using interpretable features (protein interaction, evolutionary conservation); well-established

448 pairs whose drivers fall outside our 40-gene panel, such as STAG2→STAG1 [44], reinforce the gen-
 449 erality of paralog compensation but are not evaluable in our driver-conditioned framework. Baseline
 450 essentiality features carry much of the predictive signal in dependency-based models [45, 46], consistent
 451 with our finding that the necessity component alone reaches AUROC = 0.642. Combinatorial CRISPR
 452 platforms have since mapped paralog synthetic lethals experimentally [20–23, 47, 48], converging on 13
 453 candidate cross-study gold standards [23] whose replication across screens remains imperfect [49]. Our
 454 analysis extends this prior work in several respects. We benchmarked DD against eight published deep
 455 learning methods in the CV3 setting. The seven-cohort CPTAC proteomic survey revealed protein-level
 456 co-variation patterns not examined in earlier studies. Venkatesh et al. [34] independently demonstrated
 457 widespread paralog protein compensation in CPTAC tumors driven mainly by hemizygous copy-number
 458 loss, echoing the CYCLOPS paradigm of collateral dependency from passenger deletions [50, 51]; our
 459 analysis is complementary in that we condition on driver point mutations rather than copy-number loss
 460 and couple the proteomic signal to dependency-based prioritization. We also added MSI and mutation-
 461 type stratification plus a dependency-window framework for translational prioritization. Together, the
 462 five orthogonal evidence layers go beyond a gene-family dependency atlas toward a prioritized candidate
 463 list for experimental follow-up. Two combinatorial screens published in 2025 extend this map at much
 464 larger scale: Harle et al. [24] tested 472 nominated pairs in 27 cell lines and validated 117 synthetic-lethal
 465 interactions, and Flister et al. [25] screened the near-complete paralogome (36,648 pairs) in two models,
 466 identified 146 lethal pairs, and trained predictive models on screen-derived features. Our contribution is
 467 orthogonal to both: rather than expanding the digenic map or fitting multi-feature predictors, we show
 468 that a single transparent quantity computable from existing monogenic dependency atlases recapitulates
 469 digenic signal in both screens, while making explicit that the driver-conditioned, directional component
 470 of the proposition cannot be adjudicated at current screen scale.

471 Both DD and the curated labels come from DepMap CRISPR data, which raises a circularity concern.
 472 We addressed this with several approaches. First, the headline estimate is robust to direction-unsupported
 473 labels but attenuates when DepMap-era labels are removed: relabelling the pair whose direct evidence
 474 supports only the reciprocal direction (EP300→CREBBP) as non-positive left AUROC unchanged at
 475 0.629, whereas removing the two DepMap-derived pairs reduced it to 0.613 and restricting to the eight
 476 pairs with pre-DepMap experimental evidence gave 0.525. Because the two DepMap-derived labels were
 477 themselves curated from this dataset, their direction consistency is partially expected, and the attenuation
 478 after their removal quantifies how much of the headline signed-direction signal rests on dataset-internal
 479 curation rather than pre-existing experimental evidence. Both lineage-evaluable positive entries on the

Tier A $\cup$ Tier B literature-derived benchmark (ARID1A $\rightarrow$ ARID1B in two lineages) ranked above all unlabeled controls; the remaining benchmark pairs had too few qualifying driver-mutant lines for evaluation. Second, CPTAC proteomics, PRISM drug sensitivity, and TCGA survival provide independent data layers. Third, prospective validation against SL pairs catalogued after the DepMap 26Q1 release is planned.

An important distinction separates the two external digenic-screen results. The analysis that produced statistically significant discrimination (Harle AUROC = 0.619, Flister AUROC = 0.696; both permutation $p \leq 2 \times 10^{-4}$) scored pairs by mutation-agnostic $\max|\text{DD}|$, which discards both mutation-conditioning and directionality—the two properties that define signed DD. Conversely, the genotype-stratified, directional analysis (layer 2) was uninformative: only 3 curated pairs per dataset carried any driver-mutant models (9 and 6 mutant observations), and no per-pair test reached significance after correction. The evidence that survives external scrutiny is therefore *dependency displacement*—the observation that monogenic dependency shifts recapitulate digenic SL signal—rather than *mutation-conditioned, directional compensation* per se. We regard this as concordant but circumscribed evidence: the external results demonstrate that dependency displacement carries reproducible digenic signal at scale on independent experimental platforms, but they do not validate signed DD as a mutation-conditioned prioritization tool beyond the broad concordance already established. Resolving the directional component will require genotype-powered digenic screens with deliberate enrichment for driver-mutant models, a design that neither Harle et al. nor Flister et al. undertook. Until such data exist, the directional, mutation-conditioned proposition that defines signed DD should be regarded as a design characteristic supported by internal consistency and mechanistic plausibility (e.g., the truncating-mutation gradient; PIK3CA $\rightarrow$ PIK3CB reverse-direction shifts) rather than as an externally validated performance property.

Under the within-driver multiple-testing framework (Benjamini–Hochberg correction across each driver’s HGNC paralogs), ARID1A $\rightarrow$ ARID1B does not pass the discovery threshold ($q = 0.39$ in Ovarian; $q = 0.88$ in Endometrial): ARID1A has >200 annotated paralogs, so the within-driver family is extremely large and power is limited. We therefore make no family-wise significance claim for this pair. Because ARID1A $\rightarrow$ ARID1B is itself a Tier B benchmark pair, its recovery serves as a second positive control for the framework (after SMARCA4 $\rightarrow$ SMARCA2), not as a novel discovery. Its nomination rests instead on convergent, complementary criteria: the largest lineage-level effect size among tested pairs (DD = 0.386 in Ovarian), the leading dependency-window profile among selective candidates (DWS = 2.82; mean selectivity +0.28), and strong prior biological evidence [27]. The multivariate regression reported above (association of ARID1B dependency with ARID1A mutation after TP53 and

CNV adjustment) answers a different question from the within-driver FDR test, conditional association rather than family-wise enrichment, and we do not use its low p-value to override the FDR result. ARID1A→ARID1B should thus be read as a high-priority validation candidate whose statistical evidence is incomplete, not as a statistically confirmed synthetic-lethal pair.

Because the two HIGH.SELECTIVITY pairs are both established synthetic lethals, the framework’s discovery value rests on its unlabeled candidates. The highest-ranked non-benchmark candidate is KMT2D→KMT2C (signed DWS = 1.26, MODERATE selectivity +0.056; composite score 0.525, fourth overall): KMT2D is one of the most frequently mutated epigenetic regulators across solid tumors, KMT2C and KMT2D are the catalytic paralogs of the COMPASS-like H3K4 methyltransferase complexes, and the pair shows high sequence-level similarity (0.829) with conserved domain architecture. Its dependency shift is modest ($|DD| = 0.064$), so we nominate KMT2D→KMT2C for combinatorial-CRISPR testing rather than claiming a validated dependency. This pair was also included in the unbiased digenic screen of Flister et al. [25] and was not called lethal in either of the two models screened (NCI-H1299 and MDA-MB-231, neither of which carries a KMT2D mutation), consistent with the low genotype penetrance of digenic interactions and with the modest dependency shift we observe, but an external caveat we report explicitly. RB1→RBL1 shows a reverse-direction shift, so its signed DWS is zero and it is not a selective candidate (its former $|DD|$ -based DWS of 3.44 rested on a near-zero essentiality denominator with selectivity ≈ 0.007).

1. *Sample size*: Several rare driver mutations appear in fewer than five DepMap cell lines and could not be analyzed in the primary framework; some lineages lacked sufficient samples for MSI and mutation-type stratification.
2. *Evaluation-set heterogeneity*: The evaluation set mixes functional analogs (e.g., BRCA1↔BRCA2, whose low k -mer Jaccard ≈ 0.23 reflects short 3-mers shared by chance between long proteins rather than true sequence homology) with true sequence paralogs. Stratifying by paralog type in future evaluations would improve clarity.
3. *Mutation-agnostic gold standards*: Conditioning on natural mutations of the 13 cross-platform in4mer gold-standard pairs [23], predominantly housekeeping genes whose alterations are mostly passenger events, did not separate them from unlabeled control pairs (10 of 13 pairs evaluable at the ≥ 3 -mutant threshold: AUROC = 0.50, permutation $p = 0.51$; only 3 pairs met the ≥ 5 threshold: exploratory AUROC = 0.68). The three pairs with the largest mutant-line counts (ARFGEF1/ARFGEF2, HDAC1/HDAC2, SLC25A28/SLC25A37) each exceeded the control 90th per-

centile (Supplementary Fig. S2), suggesting DD recovers signal when mutations are abundant and functionally selected. DD is therefore scoped to driver-mutation-conditioned dependency; scored against direct combinatorial assays at adequate scale, mutation-agnostic $|DD|$ does separate digenic hits from non-hits (Results), but those screens cannot test the driver-conditioned proposition itself (next item).

4. *External digenic screens:* The two combinatorial screens used as external tests contain few models carrying our driver mutations (three evaluable pairs per dataset; 9 and 6 mutant model-observations), so the genotype-stratified test is uninformative rather than confirmatory or contradictory. Both datasets were publicly released before we accessed them (2023–2024 data releases and preprint) and were analysed without modifying any frozen pipeline component; they therefore constitute independent external validation but not a prospectively registered hold-out.
5. *Circular validation:* DD and the curated labels both derive from DepMap data; the core AUROC measures internal consistency, not independent prediction. CPTAC, PRISM, and TCGA provide partially independent support but do not directly validate SL prediction.
6. *Observational design:* DD captures association, not causation. We controlled for lineage, co-mutations, and CNV through multivariate regression but cannot eliminate residual confounding. CRISPR double-knockout in isogenic models remains the necessary next step.
7. *In vitro therapeutic window:* DWS is based on cell line dependencies and cannot predict organ-specific toxicity. Integration with GTEx, IMPC, and gnomAD constraint data is needed.
8. *Structural analysis:* Limited to sequence-level features and ESMFold-predicted single-chain structures truncated to 400 residues; the reported “structural similarity” is a feature-space heuristic ($1 - |\Delta\text{feature}|$), not a three-dimensional alignment metric (no TM-score or RMSD), and the descriptors are not calibrated against drug-development data. AlphaFold-Multimer complex predictions and experimental binding data are needed for structure-guided drug design.
9. *Positive–unlabeled evaluation:* Unlabeled controls are HGNC paralog pairs without curated SL evidence, not experimentally verified non-SL pairs. AUROC estimates therefore measure separation from *unlabeled* pairs and are optimistic relative to a design with validated negatives. Under random hidden-positive contamination at rates $\pi = 1\%$, 5% , and 10% (2,000 seeded relabelings), the primary AUROC of 0.629 attenuates to means of 0.615, 0.580, and 0.557, remaining above 0.5 in $\geq 83\%$ of replicates even at $\pi = 10\%$ (Supplementary Table S13). The bias direction is unknown: adversarially placed hidden positives can drive AUROC to 0.27, while top-ranked hidden

positives raise it to 0.85; we therefore do not assert a unidirectional bias.

10. *Evaluation units*: AUROC is computed at different units across analyses: driver $\times$ paralog $\times$ lineage entries in the primary framework versus aggregated pairs in the per-pair framework; the two are not directly comparable, and pair-level aggregation discards lineage structure.
11. *Benchmark scope*: The evaluation universe comprises three gynecological lineages (Ovarian, Endometrial, Cervical); per-lineage AUROCs outside this universe (Fig. 1a) rest on few positives each and are descriptive rather than independent benchmarks.
12. *Effect direction*: The primary benchmark ranks pairs by signed DD, so reverse-direction shifts (weaker paralog dependency in driver-mutant lines) count against a pair; the direction-agnostic $|\text{DD}|$ variant, which can admit such pairs, is reported as a sensitivity analysis.
13. *Driver-mutation definitions*: TSG and oncogene mutation calls follow different rules (LikelyLoF vs. Hotspot), so “mutation” is not a uniform intervention across gene classes; amplification-driven oncogenes (e.g., CCNE1) contribute no mutant lines under these rules and are only analyzed through copy-number channels.
14. *CPTAC coverage*: Not all paralogs were quantified in all cohorts (41–45 of 58 per cohort); mutation-conditioned analysis was performed only for ARID1A $\rightarrow$ ARID1B in UCEC.

(1) Combinatorial-CRISPR validation of the nominated unlabeled candidates, starting with KMT2D $\rightarrow$ KMT2C in KMT2D-mutant versus wild-type isogenic lines, with viability and colony formation as primary read-outs; for ARID1A $\rightarrow$ ARID1B, digenic knockout epistasis in ARID1A-mutant lines (e.g., TOV-21G) would extend the single-gene evidence of Helming et al. [27] to a dual-perturbation design. This study is purely computational; experimental validation through collaboration with functional genomics laboratories is the essential next step. (2) External hold-out validation against SL pairs catalogued in SynLethDB or other databases after the DepMap 26Q1 release. (3) Extending mutation-conditioned proteomic analysis from ARID1A $\rightarrow$ ARID1B in UCEC to additional paralog pairs and CPTAC cohorts. (4) Experimental assessment of ARID1B degradation strategies, informed by the SMARCA2-degrader precedent (Discussion), once suitable ARID1B ligands or degrader handles are identified. (5) Building lineage-specific paralog-SL atlases with MSI-aware annotation to support translational research.

Conclusions

Signed DD provides a direction-aware and experimentally interpretable statistic for prioritizing mutation-conditioned paralog dependencies, evaluated here across 23 solid tumor types. Its current performance estimates remain uncertain because the literature-derived positive set is small, heterogeneous, and partly dependent on DepMap-derived evidence. As a fully external test, mutation-agnostic dependency displacement recapitulated digenic synthetic lethality signal in two independent combinatorial CRISPR screens (Harle et al. 2025; Flister et al. 2025), though the directional, mutation-conditioned proposition awaits genotype-powered screens. The framework recovered established mutation-conditioned dependencies, including ARID1A→ARID1B, the highest-selectivity established candidate, and prioritized KMT2D→KMT2C as the highest-ranked non-benchmark candidate for further testing, but neither DD nor the dependency-window score establishes causal synthetic lethality or an in vivo therapeutic window. Independent combinatorial perturbation in isogenic models is the essential next step.

We provide a complete open-source toolkit, the paralogSL R package and a reproducible Python pipeline, which together require only publicly available DepMap data and no specialized computational infrastructure.

Methods

Data sources and preprocessing

DepMap 26Q1. CRISPRGeneEffect.csv (420 MB; Chronos-normalized gene dependency scores (hereafter “gene effect scores”) for 1,208 cell lines × 18,531 genes), OmicsSomaticMutations.csv (554 MB), OmicsExpressionProteinCodingGenesTPMLogp1.csv (483 MB), and Model.csv (681 KB) were downloaded from the DepMap Portal (release 26Q1) [52, 53]. Driver-mutation status was defined by gene-class-specific rules on the default omics profile per model (IsDefaultEntryForModel) [54]: for tumor suppressors (TSG), a line is mutant if and only if it carries at least one DepMap-curated likely loss-of-function variant (LikelyLoF, which includes dominant-negative TP53 missense); for oncogenes, if and only if it carries at least one recurrent oncogenic hotspot variant (Hotspot); genes were classified as TSG or oncogene by the DepMap 26Q1 OncogeneHighImpact/TumorSuppressorHighImpact majority vote (one documented override: GATA3 as TSG, resolving a 20/20 majority-vote tie given its frameshift-dominated lesion spectrum; GATA3 drives none of the curated benchmark pairs). Amplification-driven oncogenes without qualifying hotspot mutations (e.g., CCNE1) contribute no mutant lines under this rule

and are analyzed through copy-number channels instead. The complete per-gene variant classification table (old non-silent definition vs. class-specific rule, cell line counts, and variant-type distribution) is provided as Supplementary Table S6. In the mutation-type analysis, multiple variants in the same gene per line are collapsed to the most severe consequence. Cell lines without expression or dependency data were excluded, yielding 1,208 cell lines. The exact files analysed (release filenames, sizes, SHA256 checksums, and per-file download dates) are tabulated in the repository data manifest (data/README.md), so the identical inputs can be retrieved and verified.

HGNC gene families. The HGNC complete set was downloaded from the HGNC repository [55]. Protein-coding genes with “Approved” status and non-empty gene_group_id were retained (11,050 genes). All pairwise gene group combinations yielded 66,595 paralog pairs.

CPTAC proteomics. Seven CPTAC cohorts were retrieved via the cBioPortal API [37, 38]: BRCA ($n = 122$), COAD ($n = 97$), LUAD ($n = 110$), GBM ($n = 99$), PDAC ($n = 140$), UCEC ($n = 95$), and LUSC ($n = 108$); these seven cohorts constitute the main proteomic analysis (771 tumors in total). Ovarian HGSC global proteome data (NCI Proteomic Data Commons study PDC000360, $n = 232$) [39] were downloaded for reference but were not used in the main analysis. Pearson correlations were computed on $\log_2$ -transformed abundance values (≥ 10 paired samples per test).

PRISM drug sensitivity. The PRISM Repurposing secondary screen dataset ($1,482$ compounds $\times$ 727 cell lines) was downloaded from the DepMap Portal [40]. $\Delta\text{AUC} = \text{mean}(\log_2\text{AUC}_{\text{MUT}}) - \text{mean}(\log_2\text{AUC}_{\text{WT}})$. Selective drugs: BH $q < 0.25$, $\Delta\text{AUC} < 0$, enrichment > 0.1 .

TCGA survival data. Overall-survival annotations and tumor mRNA expression for the TCGA Pan-Cancer Atlas breast cancer cohort were retrieved via the cBioPortal API (study brca_tcga_pan_can_atlas_2018) [37, 56]. Patients with both overall-survival and expression data ($n = 1,069$; 151 deaths) were analyzed. Cox proportional-hazards models (Wald 95% confidence intervals) were fitted on continuous $\log_2$ -transformed, z-scored expression, estimating hazard ratios per +1 SD for each of the 32 paralog genes from the 16 priority pairs; multivariable models adjusted for age and AJCC stage (ordinal I–IV). The proportional-hazards assumption was assessed with a Schoenfeld-type residual–time correlation test (Supplementary Table S8). The identical pipeline was extended to the UCEC ($n = 525$, 87 events) and OV ($n = 299$, 180 events) PanCan Atlas cohorts; AJCC stage annotation is defined but entirely unpopulated for these two studies in cBioPortal (verified at patient and sample level), so age-only adjusted models are the adjusted estimates of record for them (Supplementary Information, TCGA survival extension).

Mutational co-occurrence. For each priority driver–paralog pair, a 2×2 contingency table of driver-

rule mutation status (TSG: LikelyLoF; oncogene: Hotspot) was built across the 1,208 dependency-profiled DepMap 26Q1 cell lines. Co-occurrence odds ratios and two-sided p -values were computed with Fisher’s exact test; Wald 95% confidence intervals on $\log(\text{OR})$ used a 0.5 Haldane correction when any cell was zero.

Delta Dependency computation

DD was computed as defined in Equation 1, using Welch’s t -test (unequal variances); effect sizes are reported as Cohen’s d (pooled SD) and the small-sample-corrected Hedges’ g . Raw DD was chosen as the primary ranking metric over standardized alternatives (Welch t , Hedges’ g) because (i) on the primary frame the three statistics are nearly interchangeable (Spearman $\rho \geq 0.97$; paired-bootstrap ΔAUROC 95% CIs include zero; Supplementary Table S13), and (ii) on the external digenic-screen frames, raw DD is the only statistic that retains signal (Harle: 0.619 vs 0.476/0.503; Flister: 0.696 vs 0.514/0.536), because standardized statistics amplify small- n stratum noise under the max-across-lineages selection. Unless stated otherwise, DD is scored as a signed quantity (positive = compensation); analyses scored by $|\text{DD}|$, which additionally credits reverse-direction shifts, are reported as sensitivity analyses. The primary analysis required a minimum of five mutant and five wild-type cell lines per driver $\times$ lineage stratum; the more permissive $\geq 3/\geq 3$ threshold was re-run as a sensitivity analysis (Results). Benjamini-Hochberg correction [57] ($q < 0.20$) was applied within each driver gene: for driver D , all (D, P_i, c_j) combinations were corrected jointly. This is a discovery-stage pipeline; global FDR correction across all 66,595 pairs $\times$ 40 drivers $\times$ 23 cancer types was not performed because the goal was to maximize sensitivity for candidate identification. Candidates should be treated as hypothesis-generating. At a stricter $q < 0.10$, PIK3CA $\rightarrow$ PIK3CB remained significant ($q = 0.086$ in Endometrial), whereas ARID1A $\rightarrow$ ARID1B did not survive within-driver correction ($q = 0.50$), not because of weak signal ($|\text{DD}| = 0.30$) but because ARID1A has >200 HGNC paralogs.

Benchmark comparison

Curated pairs and unlabeled paralog controls. Twelve paralog-SL pairs were curated from published studies (Supplementary Table S3), which assigns each pair an evidence tier after verification of the primary citations. *Tier A* (three pairs: AKT1 $\rightarrow$ AKT2 [28], CDK4 $\rightarrow$ CDK6 and MAP2K1 $\rightarrow$ MAP2K2 [20]) comprises pairs with direct genetic synthetic-lethal evidence from dual-gene perturbation (combinatorial or digenic CRISPR knockout). *Tier B* (two pairs: SMARCA4 $\rightarrow$ SMARCA2 [26], ARID1A $\rightarrow$ ARID1B [27]) comprises pairs for which the paralog is a demonstrated selective dependency in driver-mutant cells

(natural-genotype conditional dependency from single-gene perturbation with functional validation). The Tier A $\cup$ Tier B set (five pairs) constitutes the primary literature-derived benchmark; the two tiers differ in evidence independence: Tier A derives from dual-gene perturbation, whereas Tier B is itself natural-genotype conditional dependency, the same data type DD measures, so Tier B concordance is partially expected by construction. *Tier C* (five pairs) comprises indirect evidence only and is excluded from the primary benchmark: EP300 $\rightarrow$ CREBBP, established experimentally in the reciprocal direction only (CREBBP $\rightarrow$ EP300 [29, 30]); PIK3CA $\rightarrow$ PIK3CB, whose primary evidence supports PTEN $\rightarrow$ PIK3CB [31]; CCNE1 $\rightarrow$ CCNE2, supported by developmental redundancy in mouse double knockouts [32]; and FBXW7 $\rightarrow$ FBXW2 and PPP2R1A $\rightarrow$ PPP2R1B, derived from DepMap analyses. Two mechanistic comparators (BRCA1 $\leftrightarrow$ BRCA2 [4] and STK11 $\rightarrow$ SIK1, an LKB1–SIK axis reference [58]) are not sequence paralogs and serve as specificity references only. Of the twelve pairs, 10 are true sequence paralogs (shared gene family, conserved domains), 1 is a partial homolog (STK11 $\rightarrow$ SIK1), and 1 is a functional analog (BRCA1 $\leftrightarrow$ BRCA2). Unlabeled controls comprised all remaining HGNC-defined paralog pairs involving the same 40 driver genes that did not appear in the positive set; these controls are unlabeled rather than experimentally confirmed non-SL pairs, so the evaluation is a positive–unlabeled setting and AUROC quantifies discrimination against unlabeled controls, not against verified negatives. The primary ranking score is signed DD (Equation 1); the direction-agnostic $|\text{DD}|$ variant is reported only as a sensitivity analysis because it also promotes reverse-direction effects (dependency decreasing after driver mutation, e.g., oncogene-addiction displacement for PIK3CA $\rightarrow$ PIK3CB). We audited every evaluation frame for directionality (output/direction_audit.json): on the Tier A $\cup$ Tier B benchmark set, signed and absolute scoring coincide (AUROC = 1.000), whereas on the full curated set 3 of 8 positive entries have DD < 0 (signed-DD AUROC = 0.629 vs. $|\text{DD}|$ sensitivity 0.676); full-set sensitivity values should therefore be read as direction-agnostic discrimination, and directional claims rest on signed DD. The evaluation is highly imbalanced (12 curated positives vs. $\sim$ 194 unlabeled controls per lineage). AUROC is insensitive to class prevalence, but precision–recall performance is not; we therefore report AUPRC (average precision) alongside AUROC, with bootstrap confidence intervals for both. Directional pairs (e.g., BRCA1 $\rightarrow$ BRCA2 and BRCA2 $\rightarrow$ BRCA1) are scored separately when both driver genes are in the 40-gene panel.

Contextual comparison with published methods.

DD was compared against eight published SL prediction methods evaluated in the CV3 setting by Feng et al. [15]. Published AUROCs are from the CV3 evaluation reported in Supplementary Data 1 of that study (NSMRand negative sampling, 1:1 positive:negative ratio, complete dataset); DD was

evaluated on the paralog-SL prediction task using the 12 curated paralog-SL pairs as the positive set. AUROC was computed with the rank-based Mann–Whitney statistic (identical to the scikit-learn implementation [59]); AUPRC is reported as average precision. Bootstrap 95% CIs used 10,000 resamples of evaluation entries (or pairs, for per-pair frames) with replacement; paired bootstrap resamples were used for head-to-head comparison of signed DD against component scores on identical resamples. For the DD + identity-proxy subset (3 pairs), bootstrap CI was computed separately to account for the reduced sample size. Negative controls: 100 permutations of shuffled known-SL labels.

External mutation-agnostic gold-standard check. The 13 cross-platform paralog gold standards from the in4mer digenic-knockout meta-analysis [23] were evaluated separately from the driver-conditioned benchmark because they are defined by double-knockout phenotype rather than by natural driver alterations. For each pair, both orientations (either member as the conditioning gene) were scored by the maximum |DD| across lineages, with mutation status defined by the pipeline’s gene-class rules (genes outside the 40-driver panel retain the permissive non-silent filter). Controls were 400 unlabeled paralog pairs sampled (seed 42) from the same Ensembl/HGNC pair universe, excluding the known-SL set and the in4mer pairs, and were scored identically. AUROC used the Mann–Whitney statistic with 1,000 bootstrap resamples (95% CI) and 10,000 label permutations (seed 42).

External digenic screens. Two published combinatorial CRISPR datasets were analysed as external tests with every pipeline component frozen before data access (code, features, mutation rules, and candidate list; engine fidelity against the frozen in4mer scoring: maximum absolute score difference 2.8×10^{-16} across 413 pairs). From Harle et al. [24] we used the long-format genetic-interaction table (their Table S4; 472 pairs $\times$ 27 lines) with their reported hit criterion (mean normalized genetic-interaction score < -0.5 and FDR < 0.01 , with neither single target depleted; a pair is a hit if hit in ≥ 1 line; 117/472 pairs). From Flister et al. [25] we used the per-model lethal calls (their Table S3; 146 pairs lethal in NCI-H1299 and/or MDA-MB-231) and, for genotype stratification, the 49-model interaction meta-analysis (their Table S6; hit if $\text{dzL2FC} \leq -2$, reproducing the study’s own SL counts for 98.3% of pairs). Where preprint text and tables differed on hit totals, table values were used. Each pair was scored by $\max|\text{DD}|$ across both orientations and all lineages under the ≥ 3 -mutant/ ≥ 3 -wild-type frame. Layer 1 (mutation-agnostic) compared scores between screen hits and non-hits by Mann–Whitney AUROC with 10,000 label permutations (seed 42). Layer 2 (genotype-stratified) projected the pipeline’s class-specific driver-mutation rules onto screen models via DepMap identifiers (26/27 Harle lines; 49/49 Flister models) and compared digenic hit rates (one-sided Fisher exact) and interaction scores (one-sided Mann–Whitney) between driver-mutant and wild-type models, per pair and pooled,

with Benjamini–Hochberg correction within each dataset. Analysis code: `external_holdout.py`; full outputs: `output/external_holdout.json` and companion CSV files.

Evaluation frameworks

We report DD performance under three aggregation choices, all on the same benchmark universe: the twelve curated pairs plus unlabeled paralog controls, evaluated in the three gynecological lineages (Ovarian, Endometrial, Cervical) carried through the full pipeline (110 driver $\times$ paralog $\times$ lineage entries, 8 positive entries; 72 unique pairs, 6 positive pairs). (i) The lineage-level evaluation (AUROC = 0.629; pair-clustered bootstrap 95% CI 0.253–0.933; pair-level permutation $p = 0.122$; effective sample size ≈ 85 of 110 entries) preserves lineage-specific signed-DD values and is the primary framework. (ii) The per-pair evaluation with max aggregation (AUROC = 0.598) scores each pair by its largest signed DD across the three lineages; after the class-specific driver-mutation rules shrank the evaluable positive set, it does not significantly exceed the permutation null ($p = 0.223$), a power limitation we report explicitly rather than a framework we rely on. Bootstrap and permutation analyses use this framework. (iii) The per-pair evaluation with mean aggregation (AUROC = 0.672; permutation $p = 0.090$) averages signed DD across lineages and is the aggregation used for the per-pair classifier features (Methods). Frames (ii) and (iii) differ only in the aggregation function, and neither is directly comparable to (i).

Sequence identity computation

K -mer-based Jaccard similarity ($k = 3$) was used as a computationally efficient proxy for global protein sequence identity. This metric was validated against Needleman-Wunsch global alignment identity in a subset of 50 protein pairs (13 known paralog pairs and 37 non-paralog controls; Pearson $r = 0.88$, $p < 10^{-16}$; Supplementary Fig. S9). Within-group correlations were $r = 0.91$ (paralogs, $n = 13$) and $r = 0.41$ (non-paralogs, $n = 37$). Within the $\geq 30\%$ identity range where the proxy is used for filtering, the RMSE was 0.12 (mean absolute error 0.09, $n = 8$ pairs with both measurements), so the $\geq 30\%$ threshold should be interpreted as an approximate enrichment filter rather than a precise identity cutoff. The $\geq 30\%$ threshold should therefore be interpreted as an approximate filter for enriching high-identity pairs rather than a precise sequence identity cutoff.

Potential confounders and causal interpretation

DD quantifies an observational association between driver mutation status and paralog dependency, not a causal effect. Potential confounders include: (1) lineage differences between mutant and wild-type

cell lines within a cancer type; (2) co-occurring mutations that may independently affect paralog dependency; (3) copy number alterations at paralog loci. We mitigated these through: within-lineage analysis (DD computed separately per cancer type), Welch's *t*-test (robust to unequal variances), and CNV independence verification (Supplementary Fig. S3, all $R^2 < 0.10$). Benjamini-Hochberg correction was applied within each driver gene to control false discovery rate. Nevertheless, experimental validation (e.g., CRISPR double-knockout in isogenic cell lines) remains essential to establish causality.

CPTAC proteomic co-variation

Protein abundance data for seven CPTAC cohorts (BRCA, COAD, LUAD, GBM, PDAC, UCEC, LUSC; $n = 771$ total) were retrieved via the cBioPortal API [37, 38]. For each paralog pair, Pearson correlation of protein $\log_2$ -abundance was computed across all samples within each cohort. Significance was assessed using BH correction within each pair across 7 cohorts. For the RNA-layer comparison, continuous mRNA profiles (`rna_seq_v2_mrna`, log-scale values used as provided) were retrieved for the same 26 pairs in all seven cohorts (179 pair $\times$ cohort tests with ≥ 10 samples), and protein-mRNA comparisons used the 122 pair $\times$ cohort combinations with both measurements. Tumor-purity robustness was assessed by partial Pearson correlation controlling for the cohort's purity annotation (ESTIMATE tumor purity in BRCA, LUAD, and LUSC; CIBERSORT absolute score in UCEC; 69 tests with ≥ 15 matched samples). Mutation-conditioned proteomic analysis was performed for the ARID1A $\rightarrow$ ARID1B pair in the UCEC cohort: ARID1A mutation status was retrieved from the cBioPortal mutations API (`ucec_cptac_2020_mutations`), and ARID1B protein abundance was compared between ARID1A-mutant ($n = 37$) and wild-type ($n = 58$) samples using the Wilcoxon rank-sum test. Extension to additional pairs and cohorts is a priority future direction.

Pan-cancer lineage definitions

Twenty-seven solid tumor types were defined by matching OncotreePrimaryDisease annotations to curated keyword lists; the 23 types with ≥ 6 cell lines with available data were analyzed (Supplementary Table S1). Lineage short labels follow DepMap/Oncotree conventions and mix organ-based and histology-based names as required by each cancer type's classification (e.g., NSCLC and SCLC are distinct histological subtypes of lung cancer); full names are given in Supplementary Table S1. Forty pan-cancer driver genes were curated from TCGA PanCanAtlas driver analyses [35] and the COSMIC Cancer Gene Census [36] (Supplementary Table S4); per-gene mutant-line counts under the final class-specific rules are reported in Supplementary Table S6 (CCNE1 contributes zero mutant lines under the

815 hotspot rule; MAPK1, audited for the external in4mer analysis rather than as a panel gene, likewise
816 yields zero).

817 **MSI stratification**

818 Cell lines were classified with the official DepMap 26Q1 microsatellite annotation (`OmicsglobalSignatures.csv`).
819 The DepMap consortium computes an MSIsensor2 MSIScore from each model's sequencing data; fol-
820 lowing the consortium convention, lines with MSIScore > 20 were labelled MSI-H and all others MSS.
821 Default-entry profiles (`IsDefaultEntryForModel`) were used; all 87 evaluable colorectal and en-
822 dometrial lines had official annotations. Because the annotation is measured rather than mutation-
823 inferred, POLE-mutant lines receive their observed status (15 of 22 POLE-mutant lines in these two
824 lineages are officially MSI-H) and are not presumed hypermutated; agreement between this annota-
825 tion and a damaging-mutation proxy (MLH1/MSH2/MSH6/PMS2/POLE) was 89% (per-line audit at
826 `output/msi_classification_audit.csv`). DD was computed separately within each subgroup for
827 colorectal and endometrial cancers. For the higher-powered interaction analysis, per-pair ordinary least-
828 squares models were fitted on all lineage lines with dependency data and official MSI annotation (col-
829 orectal $n = 63$, endometrial $n = 28$): $G_P \sim D_{\text{mutation}} \times \text{MSI-H}$, with driver mutation status from the
830 gene-class rules described above; the interaction coefficient equals the MSI-H minus MSS difference in
831 DD (ΔDD). A pair was evaluable with ≥ 3 mutant lines and ≥ 2 mutants in each MSI class (47 pairs
832 in colorectal, 24 in endometrial). A continuous variant replaced the MSI-H indicator with $\log(1 +$
833 `MSIsensor2` score). Interaction p -values were combined per lineage with Stouffer's method (signed
834 by ΔDD), the excess of nominally significant interactions was assessed with a one-sided binomial test
835 against the 5% null rate, and Benjamini–Hochberg q -values were computed within each lineage (`msi_interaction_test`
836 per-pair results at `output/msi_interaction_results.csv`).

837 **Mutation type classification**

838 Driver mutations classified as truncating (frameshift, nonsense, splice-acceptor, splice-donor, start-loss,
839 stop-loss) or missense (`missense_variant`, `protein_altering_variant`, excluding truncating) using VEP Vari-
840 antInfo annotations. Per-cell-line classification used the most severe mutation type.

841 **Therapeutic window analysis**

842 DWS defined as in Equation 2 with the signed numerator $\max(DD, 0)$, which credits only compensation-
843 direction shifts; the $|DD|$ -numerator variant is reported as a sensitivity analysis (Supplementary Ta-

ble S10). The denominator uses $\max(|\mu_{\text{Chronos}}|, f_{\text{pan-essential}}, 0.01)$ to capture the most conservative estimate of baseline paralog essentiality: $|\mu_{\text{Chronos}}|$ reflects mean dependency magnitude (dimensionless, normalized by Chronos), $f_{\text{pan-essential}}$ reflects the fraction of lines where the paralog is essential (dimensionless, range 0–1), and the 0.01 floor prevents division by near-zero values for non-essential paralogs. All three quantities are dimensionless ($f_{\text{pan-essential}}$ is bounded in 0–1; $|\mu_{\text{Chronos}}|$ is typically of order 0–1); the max operator selects whichever indicates the greatest baseline toxicity risk. Paralogs were classified with operational stratification thresholds: PAN_ESSENTIAL ($f_{\text{pan-essential}} > 0.5$), HIGH_SELECTIVITY (selectivity > 0.15 and DWS > 1.0), MODERATE, or LOW_SELECTIVITY, which are not validated clinical cutoffs (threshold sensitivity in Supplementary Table S10). Analysis across 5 cancer contexts (Ovarian, Endometrial, Breast, Colorectal, PanCancer). The module covers 24 candidate pairs (13 curated-direction pairs, including the two comparators and the BRCA2→BRCA1 reverse orientation, plus 11 pipeline-nominated pairs); 21 had at least one context passing the ≥ 5 -mutant frame and are reported, while AKT1→AKT2, CCNE1→CCNE2, and CDK4→CDK6 did not meet the threshold in any context. Robustness: pair rankings were re-derived under alternative denominators ($|\mu_{\text{Chronos}}|$ only; $f_{\text{pan-essential}}$ only; the mean of the two), alternative floors (0.001, 0.05), and alternative selectivity thresholds (0.10, 0.20), and bootstrap 95% CIs for mean DWS and selectivity were estimated from 1,000 stratified resamples of cell lines (seed 42; Supplementary Tables S5 and S10; dws_robustness.py).

Exploratory sequence- and structure-derived targetability descriptors

Sequence-based features (length, molecular weight, pI, GRAVY hydrophobicity, disorder propensity, lysine/cysteine density) were computed from UniProt canonical sequences, and domain architectures from UniProt annotations. Structures were predicted with ESMFold v1 (facebook/esmfold_v1) for the N-terminal 400 residues of 15 proteins (mean pLDDT = 0.724); for large proteins this truncation covers only a fraction of the chain (e.g., $\sim 17\%$ for ARID1A/ARID1B), so ARID1A/ARID1B were assessed with sequence-level features only. We emphasize the exploratory status of these descriptors. (i) The reported “structural similarity” is a feature-space heuristic, $\max(0, 1 - |\Delta_{\text{feature}}|)$ over the sequence features above; it is not a three-dimensional alignment metric (no TM-score or RMSD is used in the score). (ii) The “pocket count” is a custom sliding-window heuristic (windows with pLDDT > 0.7 and hydrophobic fraction 0.2–0.7, aggregated per 12 windows), not a validated pocket-detection method. (iii) The PROTAC-suitability score (lysine density $\times 2.0 +$ disorder propensity $-$ cysteine density $\times 0.5$) captures only two of the many determinants of degrader efficacy; ligandability, E3-ligase compatibility, ternary-complex geometry, and cell permeability are not modelled. A per-pair

targetability score combined structural similarity (0.25), domain-architecture conservation (0.30), the druggability proxy (0.25), and the PROTAC score (0.20). The composite prioritization score combined max-normalized DWS (0.40), selectivity rescaled to $[0, 1]$ as $(s + 1)/2$ (0.30), and targetability (0.30); all weights were fixed heuristically before final reporting and are unchanged in the archived release (commit 880cb9f; `alphafold_analysis.py`), and none of the descriptors was calibrated against external drug-development data.

Statistical analyses

AUROC via scikit-learn [59]. Bootstrap CI: 2.5th and 97.5th percentiles from 1,000 iterations. For the primary lineage-level frame, a pair-clustered bootstrap (10,000 iterations, resampling whole pairs together with all their lineage entries) and a pair-level label permutation test (10,000 permutations) account for within-pair correlation (one-way random-effects ANOVA ICC = 0.554 for signed DD within pairs; effective sample size 85 of 110 entries). A driver-block bootstrap (10,000 iterations, resampling driver genes together with all their pairs and lineage entries) further accounts for non-independence among pairs sharing a driver gene. Negative controls: 10,000 label permutations. Analyses in Python 3 with numpy [60], scipy [61], pandas, statsmodels, and scikit-learn [59]; exact dependency versions are pinned in the repository Dockerfile; the complete audit suite passes under the current managed runtime (Python 3.12, scikit-learn 1.9.0, numpy 2.4.4, scipy 1.18.0, pandas 3.0.2, statsmodels 0.14.6). The companion R package `paralogSL` provides DD computation, PCS analysis, benchmark comparison, and CPTAC correlation visualization under the MIT license.

Multiple testing correction

Multiple testing correction was applied with Benjamini–Hochberg (BH) FDR control, but the correction unit varies by analysis because the tested hypotheses differ in scope and dependence structure. For DD computation, BH was applied within each driver gene across all of that driver’s HGNC paralogs and lineages jointly ($q < 0.20$, discovery-stage; $q < 0.10$ sensitivity analysis); this within-driver correction is a discovery-stage device and does not constitute genome-wide FDR control. For MSI interaction analyses, BH was applied within each lineage separately. For CPTAC protein correlations, BH was applied within each paralog pair across the 7 cohorts. For PRISM drug sensitivity, a deliberately permissive discovery threshold ($q < 0.25$) was used. For TCGA survival analyses, BH was applied across the 32-gene family within each model class (univariable and multivariable). Because the correction families differ, q -values from different analyses are not directly comparable, and the term “significant” is reserved for results

passing the stated FDR threshold within the stated family; all other results are described as nominal. Because the correction unit differs across analyses, q-values computed within different families are not directly comparable; cross-analysis convergence arguments therefore rest on effect sizes and consistency of direction, not on comparing q-value magnitudes.

Power analysis and test set stratification

Curated pairs were stratified into evidence tiers (Supplementary Table S3): *Tier A* (3 pairs; direct dual-perturbation genetic synthetic-lethal evidence), *Tier B* (2 pairs; natural-genotype conditional dependency with functional validation), *Tier C* (5 pairs; indirect, reciprocal-direction-only, or DepMap-derived evidence), and 2 mechanistic comparators (BRCA1 $\leftrightarrow$ BRCA2, STK11 $\rightarrow$ SIK1). The primary endpoint is the AUROC and AUPRC of signed DD on the Tier A $\cup$ Tier B literature-derived benchmark at the driver $\times$ paralog $\times$ lineage level, the natural unit at which DD is computed, with the pair treated as the grouping unit in sensitivity analyses to avoid cross-lineage information leakage. The full 12-pair curated set, tier-restricted, direction-strict, and per-pair aggregated analyses are reported as secondary and sensitivity analyses. With at most 2–5 evaluable positive pairs, subset-level AUROC has limited power: at $\alpha = 0.05$, an estimated ≥ 25 validated positive pairs would be needed to achieve 80% power for detecting AUROC = 0.85 against a null of 0.5. At the observed primary AUROC of 0.629, power with the current 8 positive and 102 control entries is only 0.34 (0.28 for the 6-positive, 66-control pair frame), and 80% power at this effect size would require ≥ 45 positive pairs (binormal model, Hanley–McNeil variance; `power_analysis.py`); because these calculations assume independent entries and verified true negatives, they serve as rough planning values under the positive–unlabeled design rather than exact requirements. To compensate for this limitation, we complement binary classification metrics with effect-size ranking (DD, DWS, selectivity) and multivariate regression significance (p -values for individual pairs).

Supplementary information

Supplementary Figures S1–S10 and Supplementary Tables S1–S12 are available as a separate PDF file (Additional file 1); machine-readable TSV mirrors of the data tables are provided as Additional files 2–12 (see Additional files). Supplementary Fig. S1: Evaluation overview across 23 solid tumor types (cell line counts, tested pairs, and per-lineage DD AUROC values, colored by AUROC tier). Supplementary Fig. S2: Evaluation robustness: bootstrap distribution, label-permutation negative control, and external check against mutation-agnostic in4mer gold standards. Supplementary Fig. S3: CNV independence analysis. Supplementary Fig. S4: CPTAC per-cohort scatter plots. Supplementary Fig. S5: Full

mutation-type stratification. Supplementary Fig. S6: Mutational co-occurrence of paralog pairs with
 and without TMB adjustment, pan-cancer and per-lineage. Supplementary Fig. S7: PRISM drug selec-
 tivity heatmap of the top associations passing the discovery-stage threshold. Supplementary Fig. S8:
 Dependency-window score (DWS) ranking, selectivity-tier composition, and cross-context consistency.
 Supplementary Fig. S9: Exploratory sequence- and structure-derived descriptors: k -mer Jaccard vs.
 Needleman-Wunsch alignment identity validation ($n = 50$ pairs, Pearson r shown), structure-based drug-
 gability assessment of 15 proteins using ESMFold-predicted structures, feature-space structural simi-
 larity, domain conservation, and composite prioritization score. Supplementary Fig. S10: Paralog ex-
 pression and overall survival in TCGA BRCA (forest plot, multivariable Cox models). Supplementary
 Table S1: Cell line counts per cancer type. Supplementary Table S2: Complete paralog-SL pair analysis
 results (110 driver $\times$ paralog $\times$ lineage associations in the three evaluation lineages (Ovarian, Endometrial,
 Cervical) passing the minimum sample size filter of ≥ 5 mutant and ≥ 5 wild-type cell lines). Supple-
 mentary Table S3: Curated paralog-SL pair evidence with tier assignment. Supplementary Table S4:
 Complete 40-gene driver panel with cancer types, selection source, and functional category. Supple-
 mentary Table S5: All 21 paralog pairs ranked by signed dependency window score with numerator,
 denominator, bootstrap confidence intervals, and bootstrap rank intervals. Supplementary Table S6:
 Gene-class-specific driver-mutation rules and per-gene variant classification (TSG: LikelyLoF; onco-
 gene: Hotspot). Supplementary Table S7: Per-association effect sizes (Cohen's d , Hedges' g) and Welch
 t -test p -values for all 110 evaluated associations (TSV file). Supplementary Table S8: TCGA BRCA
 overall-survival associations for the 32 paralog genes (univariable and multivariable Cox models, FDR,
 PH tests). Supplementary Table S9: Top paralog-SL candidates ranked by the exploratory composite
 prioritization score. Supplementary Table S10: Sensitivity of dependency-window score rankings and
 selectivity classification. Supplementary Table S11: Full regression model table for confounder controls
 (base, CNV-, expression-, and lineage-adjusted models with robust standard errors and BH q -values).
 Supplementary Table S12: Contextual literature values under non-comparable evaluation settings (pub-
 lished CV3 AUROCs from Feng et al. with the DD value from this study; evaluation frameworks differ
 and direct comparison is not warranted).

Abbreviations

SL: synthetic lethality; DD: Delta Dependency; PCS: paralog compensation score; DWS: dependency
 window score; TSG: tumor suppressor gene; LoF: loss of function; CPTAC: Clinical Proteomic Tumor

Analysis Consortium; TCGA: The Cancer Genome Atlas; PRISM: Profiling Relative Inhibition Simultaneously in Mixtures; DepMap: Cancer Dependency Map; AUROC: area under the receiver operating characteristic curve; AUPRC: area under the precision–recall curve; MSI-H: microsatellite instability-high; MSS: microsatellite stable; MMR: mismatch repair; CNV: copy number variation; HNSCC: head and neck squamous cell carcinoma; SCLC: small cell lung cancer; NSCLC: non-small cell lung cancer; OR: odds ratio; HR: hazard ratio; CI: confidence interval; LR: logistic regression; RF: random forest; SVM: support vector machine.

Declarations

Ethics approval and consent to participate

No institutional ethics approval was required: this study analysed exclusively publicly available, de-identified data (DepMap, CPTAC, TCGA, PRISM).

Consent for publication

Not applicable.

Data availability

Raw third-party datasets. DepMap 26Q1 from the DepMap Portal [52] (exact release filenames, sizes, SHA256 checksums, and per-file download dates in the repository data manifest, data/README.md); CPTAC proteomics from cBioPortal [37, 38] and the Proteomic Data Commons [62]; TCGA PanCan from UCSC Xena [63]; PRISM from the DepMap PRISM portal [40]; HGNC gene families from the HGNC repository [55].

Processed analysis-ready data. Cell line-level DD values and all derived analysis tables are archived at Zenodo (concept DOI <https://doi.org/10.5281/zenodo.21502030>, which resolves to the latest release) and in the GitHub repository below.

Machine-readable supplementary tables. TSV mirrors of Supplementary Tables S1–S12 are provided as Additional files 2–12, each documented in Additional file 1.

Code. See Code availability below.

Archived manuscript-specific release. The submission snapshot is GitHub tag v1.4.2 with version DOI <https://doi.org/10.5281/zenodo.21996173> under the concept DOI above; the accepted-version snapshot will be deposited identically.

Code availability

All analysis code is available under the MIT license at <https://github.com/tjogzt/paralogSL> (concept DOI [10.5281/zenodo.21502113](https://doi.org/10.5281/zenodo.21502113); submission snapshot v1.2.0, version DOI [10.5281/zenodo.21996281](https://doi.org/10.5281/zenodo.21996281)).

The complete analysis pipeline (DOI: [10.5281/zenodo.21502030](https://doi.org/10.5281/zenodo.21502030)) is available at <https://github.com/tjogzt/paralog-sl-predictor>. The paralogSL R package (v1.2.0, 21 exported functions) provides DD computation, benchmark comparison, dependency-window scoring, CPTAC visualization, and candidate-prioritization utilities (`visualize_dependency_shift` and related helpers). Random seeds (`set.seed(42)` in R, `random_state=42` in Python) are fixed throughout.

Workflow: DepMap gene-effect matrix $\rightarrow$ `build_mutation_matrix()` $\rightarrow$ `compute_dd()` per driver $\times$ paralog $\times$ lineage $\rightarrow$ `compute_auroc()` against the curated positive set $\rightarrow$ `compute_therapeutic_window()` (historical function name retained in the package; computes the DWS) for translational ranking $\rightarrow$ visualization via `plot_pancancer_summary()` and `plot_dws()` (`plot_therapeutic_window()` is a soft-deprecated alias).

Quick start (3 lines):

```
remotes::install_github("tjogzt/paralogSL") # source: github.com/tjogzt/paralogSL
library(paralogSL)
result <- run_paralog_analysis(dep_matrix, mut_matrix, driver="ARID1A", paralog="ARID1B")
```

The package includes built-in datasets (`known_sl_pairs`, `gyn_drivers`, `solid_tumor_summary`), a vignette with worked examples, and 20 exported functions. A frozen release with all processed data tables is archived at Zenodo (concept DOI <https://doi.org/10.5281/zenodo.21502030>; submission snapshot: GitHub tag v1.4.0, commit 880cb9f, version DOI <https://doi.org/10.5281/zenodo.21765256>). The R package is archived separately (concept DOI <https://doi.org/10.5281/zenodo.21502113>; v1.1.2 version DOI <https://doi.org/10.5281/zenodo.21634707>). Each new release mints a new version DOI under the same concept DOI, and the accepted-version snapshot will be deposited identically. The complete Python analysis pipeline (`main.py`, `pcs.py`, `pancancer.py`) is included in the same repository for full reproducibility. In addition, four audit scripts (`compute_headline_metrics.py`, `ml_benchmark.py`, `regression_controls.py`, and `audit_manuscript_numbers.py`) recompute every quantitative claim in this manuscript directly from raw DepMap data and derived artifacts: headline metrics, classifier benchmarks, covariate-adjusted regressions, and all remaining numbers (lineage-level AUROCs, mechanism contrasts, MSI and mutation-type stratifications, direction audit, therapeutic-window module, pair-clustered bootstrap and permutation uncertainty estimates, power analysis, com-

posite fold-internal scaling, composite-weight and in4mer seed sensitivities, the external digenic-screen hold-out against the Harle 2025 and Flister 2025 combinatorial screens, CPTAC RNA-layer and purity-adjusted co-variation, the TCGA UCEC/OV survival extension, and machine-readable table mirrors), 420 claims in total. Each script ends with an automated claim-by-claim comparison against the reported values; a single entry point (`verify_all.sh`) runs all four audits plus the test suite in ~30 s on a standard workstation.

Competing interests

The authors declare no competing interests.

Funding

This work was supported by the National Natural Science Foundation of China (grant nos. 82403759 to C.X. and 82403616 to Y.W.).

Authors' contributions

Q.M. and T.Z. conceived the study. T.Z. developed the methodology, performed all computational analyses, developed the R package, and wrote the original draft of the manuscript. Q.M. curated clinical datasets, provided gynecological oncology domain expertise, and contributed to manuscript writing and revision. P.C., C.X., Y.W., T.H. and Q.S. contributed clinical data curation and clinical interpretation; C.X. and Y.W. additionally acquired funding. All authors read and approved the final manuscript.

Acknowledgements

We thank the DepMap, CPTAC, TCGA, PRISM, and cBioPortal teams for making their data publicly available. We also acknowledge the SynLethDB, HGNC, and UniProt teams for curated resources. During the preparation of this work, the authors used Kimi (Moonshot AI) for language editing, citation cross-checking, and code-refactoring assistance. All AI-generated suggestions were reviewed, tested, and approved by the authors, who take full responsibility for the integrity and content of this publication.

Additional files

Additional file 1: Supplementary Information (PDF). Supplementary Methods, Supplementary Figures S1–S10, and Supplementary Tables S1–S13.

Additional file 2: TableS1_CellLineCounts.tsv. Cell line counts per cancer type (23 solid tumor types).

Additional file 3: TableS2_FullResults.tsv. Complete paralog-SL pair analysis results (110 driver $\times$ paralog $\times$ cancer type associations $\times$ 17 columns).

Additional file 4: TableS3_GoldStandard.tsv. Machine-readable mirror of the 12 evidence-tiered curated pairs (file name retained for continuity).

Additional file 5: TableS4_DriverGenes.tsv. Complete 40-gene driver panel with cancer types, selection source, and functional category.

Additional file 6: TableS5_DWS.tsv. All 21 paralog pairs ranked by signed dependency window score with numerator, denominator, floor flag, and bootstrap 95% confidence intervals.

Additional file 7: TableS6_MutationRules.tsv. Gene-class-specific driver-mutation rules and per-gene variant classification (TSG: LikelyLoF; oncogene: Hotspot).

Additional file 8: TableS7_EffectSizes.tsv. Per-association effect sizes (Cohen's d , Hedges' g) and Welch t -test p -values for all 110 evaluated associations.

Additional file 9: TableS8_BRCA_Survival.tsv. TCGA BRCA overall-survival associations for the 32 paralog genes (univariable and multivariable Cox models, FDR, PH tests).

Additional file 10: TableS9_CompositeScore.tsv. Paralog-SL candidates ranked by the exploratory composite prioritization score, annotated with the curated evidence tier of each pair (Table S3).

Additional file 11: TableS10_DWS_Sensitivity.tsv. Sensitivity of dependency-window score rankings and selectivity classification.

Additional file 12: TableS11_RegressionModels.tsv. Full regression model table for confounder controls: base, CNV-adjusted, expression-adjusted, and lineage-adjusted models per driver-paralog pair, with mutation coefficients, HC3 robust standard errors, 95% confidence intervals, nominal p -values, and Benjamini-Hochberg q -values.

References

[1] O'Neil NJ, Bailey ML, Hieter P. Synthetic lethality and cancer. *Nat Rev Genet.* 2017;18(10):613–623.

[2] Huang A, Garraway LA, Ashworth A, Weber B. Synthetic lethality as an engine for cancer drug target discovery. *Nat Rev Drug Discov.* 2020;19(1):23–38.

1079 [3] Lord CJ, Ashworth A. PARP inhibitors: synthetic lethality in the clinic. *Science*.
1080 2017;355(6330):1152–1158.

1081 [4] Bryant HE, Schultz N, Thomas HD, Parker KM, Flower D, Lopez E, et al. Specific
1082 killing of BRCA2-deficient tumours with inhibitors of poly(ADP-ribose) polymerase. *Nature*.
1083 2005;434(7035):913–917.

1084 [5] Farmer H, McCabe N, Lord CJ, Tutt ANJ, Johnson DA, Richardson TB, et al. Targeting the DNA
1085 repair defect in BRCA mutant cells as a therapeutic strategy. *Nature*. 2005;434(7035):917–921.

1086 [6] Gonçalves E, Ryan CJ, Adams DJ. Synthetic lethality in cancer drug discovery: challenges and
1087 opportunities. *Nat Rev Drug Discov*. 2026;25(1):22–38.

1088 [7] Guo J, Liu H, Zheng J. SynLethDB: synthetic lethality database toward discovery of selective and
1089 sensitive anticancer drug targets. *Nucleic Acids Res*. 2016;44(D1):D1011–D1017.

1090 [8] Wang J, Wu M, Huang X, Wang L, Zhang S, Liu H, et al. SynLethDB 2.0: a web-based
1091 knowledge graph database on synthetic lethality for novel anticancer drug discovery. *Database*.
1092 2022;2022:baac030.

1093 [9] Zhu Y, Zhou Y, Liu Y, Wang X, Li J. SLGNN: synthetic lethality prediction in human cancers based
1094 on factor-aware knowledge graph neural network. *Bioinformatics*. 2023;39(2):btad015.

1095 [10] Long Y, Wu M, Liu Y, Zheng J, Kwok CK, Luo J, et al. Graph contextualized attention network for
1096 predicting synthetic lethality in human cancers. *Bioinformatics*. 2021;37(16):2432–2440.

1097 [11] Cai R, Chen X, Fang Y, Wu M, Hao Y. Dual-dropout graph convolutional network for predicting
1098 synthetic lethality in human cancers. *Bioinformatics*. 2020;36(16):4458–4465.

1099 [12] Huang J, Wu M, Lu F, Ou-Yang L, Zhu Z. Predicting synthetic lethal interactions in human
1100 cancers using graph regularized self-representative matrix factorization. *BMC Bioinformatics*.
1101 2019;20(Suppl 19):657.

1102 [13] Wang S, Xu F, Li Y, Wang J, Zhang K, Liu Y, et al. KG4SL: knowledge graph neural network for
1103 synthetic lethality prediction in human cancers. *Bioinformatics*. 2021;37(Suppl 1):i418–i425.

1104 [14] Wang S, Feng Y, Liu X, Liu Y, Wu M, Zheng J. NSF4SL: negative-sample-free contrastive learning
1105 for ranking synthetic lethal partner genes in human cancers. *Bioinformatics*. 2022;38(Suppl 2):ii13–
1106 ii19.

- 1107 [15] Feng Y, Long Y, Wang H, Ouyang Y, Li Q, Wu M, et al. Benchmarking machine learning methods
1108 for synthetic lethality prediction in cancer. *Nat Commun.* 2024;15:9058.
- 1109 [16] Hao Z, Wu D, Fang Y, Wu M, Cai R, Li X. Prediction of synthetic lethal interactions in human
1110 cancers using multi-view graph auto-encoder. *IEEE J Biomed Health Inform.* 2021;25(10):4041–
1111 4051.
- 1112 [17] Koonin EV. Orthologs, paralogs, and evolutionary genomics. *Annu Rev Genet.* 2005;39:309–338.
- 1113 [18] D’Antonio M, Guerra RF, Cereda M, Marchesi S, Montani F, Nicassio F, et al. Recessive
1114 cancer genes engage in negative genetic interactions with their functional paralogs. *Cell Rep.*
1115 2013;5(6):1519–1526.
- 1116 [19] De Kegel B, Quinn N, Thompson NA, Adams DJ, Ryan CJ. Comprehensive prediction of robust
1117 synthetic lethality between paralog pairs in cancer cell lines. *Cell Syst.* 2021;12(12):1144–1159.e6.
- 1118 [20] Parrish PCR, Thomas JD, Gabel AM, Kamlapurkar S, Bradley RK, Berger AH. Discov-
1119 ery of synthetic lethal and tumor suppressor paralog pairs in the human genome. *Cell Rep.*
1120 2021;36(9):109597.
- 1121 [21] Thompson NA, Ranzani M, van der Weyden L, Iyer V, Offord V, Droop A, et al. Combinatorial
1122 CRISPR screen identifies fitness effects of gene paralogues. *Nat Commun.* 2021;12(1):1302.
- 1123 [22] Dede M, McLaughlin M, Kim E, Hart T. Multiplex enCas12a screens detect functional buffer-
1124 ing among paralogs otherwise masked in monogenic Cas9 knockout screens. *Genome Biol.*
1125 2020;21(1):262.
- 1126 [23] Esmaeili Anvar N, Lin C, Ma X, Wilson LL, Steger R, Sangree AK, Colic M, Wang SH, Doench JG,
1127 Hart T. Efficient gene knockout and genetic interaction screening using the in4mer CRISPR/Cas12a
1128 multiplex knockout platform. *Nat Commun.* 2024;15(1):3577.
- 1129 [24] Harle V, Offord V, Gökbağ B, Fotopoulos L, Williams T, Alexander D, Mehta I, Thompson NA,
1130 Olvera-León R, Peidli S, Iyer V, Gonçalves E, Kebabci N, De Kegel B, van de Haar J, Li L, Ryan
1131 CJ, Adams DJ. A compendium of synthetic lethal gene pairs defined by extensive combinatorial
1132 pan-cancer CRISPR screening. *Genome Biol.* 2025;26:284.
- 1133 [25] Flister MJ, Verduzco D, Petiwala S, Carneiro-Lobo TC, Gekas C, Shi X, Lu C, Dezso Z. Genetic
1134 mapping and predictive modeling of paralog synthetic lethality. *Cell Rep.* 2025;44(11):116512.

- 1135 [26] Hoffman GR, Rahal R, Buxton F, Xiang K, McAllister G, Frias E, et al. Functional epigenetics
1136 approach identifies BRM/SMARCA2 as a critical synthetic lethal target in BRG1-deficient cancers.
1137 *Proc Natl Acad Sci USA*. 2014;111(8):3128–3133.
- 1138 [27] Helming KC, Wang X, Wilson BG, Vazquez F, Haswell JR, Manchester HE, et al. ARID1B is a
1139 specific vulnerability in ARID1A-mutant cancers. *Nat Med*. 2014;20(3):251–254.
- 1140 [28] Najm FJ, Strand C, Donovan KF, Hegde M, Sanson KR, Vaimberg EW, et al. Orthologous CRISPR-
1141 Cas9 enzymes for combinatorial genetic screens. *Nat Biotechnol*. 2018;36(2):179–189.
- 1142 [29] Ogiwara H, Sasaki M, Mitachi T, Oike T, Higuchi S, Tominaga Y, et al. Targeting p300 addiction in
1143 CBP-deficient cancers causes synthetic lethality by apoptotic cell death due to abrogation of MYC
1144 expression. *Cancer Discov*. 2016;6(4):430–445.
- 1145 [30] Nie M, Du L, Ren W, Joung J, Ye X, Shi X, et al. Genome-wide CRISPR screens reveal synthetic
1146 lethal interaction between CREBBP and EP300 in diffuse large B-cell lymphoma. *Cell Death Dis*.
1147 2021;12(5):419.
- 1148 [31] Wee S, Wiederschain D, Maira SM, Loo A, Miller C, deBeaumont R, et al. PTEN-deficient cancers
1149 depend on PIK3CB. *Proc Natl Acad Sci USA*. 2008;105(35):13057–13062.
- 1150 [32] Geng Y, Yu Q, Sicinska E, Das M, Schneider JE, Bhattacharya S, et al. Cyclin E ablation in the
1151 mouse. *Cell*. 2003;114(4):431–443.
- 1152 [33] Jia S, Liu Z, Zhang S, Liu P, Zhang L, Lee SH, et al. Essential roles of PI(3)K-p110 β in cell growth,
1153 metabolism and tumorigenesis. *Nature*. 2008;454(7205):776–779.
- 1154 [34] Venkatesh A, Quinn N, Upadhyaya SR, De Kegel B, Bolado Carrancio A, Lefeuvre T, et al. Proteomic
1155 compensation by paralogs preserves protein interaction networks after gene loss in cancer. *Mol Syst*
1156 *Biol*. 2025;21(8):1090–1118.
- 1157 [35] Bailey MH, Tokheim C, Porta-Pardo E, Sengupta S, Bertrand D, Weerasinghe A, et al. Compre-
1158 hensive characterization of cancer driver genes and mutations. *Cell*. 2018;173(2):371–385.e18.
- 1159 [36] Sondka Z, Bamford S, Cole CG, Ward SA, Dunham I, Forbes SA. The COSMIC Cancer Gene Cen-
1160 sus: describing genetic dysfunction across all human cancers. *Nat Rev Cancer*. 2018;18(11):696–
1161 705.

- [37] Cerami E, Gao J, Dogrusoz U, Gross BE, Sumer SO, Aksoy BA, et al. The cBio cancer genomics portal: an open platform for exploring multidimensional cancer genomics data. *Cancer Discov.* 2012;2(5):401–404.
- [38] Gao J, Aksoy BA, Dogrusoz U, Dresdner G, Gross B, Sumer SO, et al. Integrative analysis of complex cancer genomics and clinical profiles using the cBioPortal. *Sci Signal.* 2013;6(269):p11.
- [39] Edwards NJ, Oberti M, Thangudu RR, Cai S, McGarvey PB, Jacob S, et al. The CPTAC data portal: a resource for cancer proteomics research. *J Proteome Res.* 2015;14(6):2707–2713.
- [40] Corsello SM, Nagari RT, Spangler RD, Rossen J, Kocak M, Bryan JG, et al. Discovering the anticancer potential of non-oncology drugs by systematic viability profiling. *Nat Cancer.* 2020;1(2):235–248.
- [41] Farnaby W, Koegl M, Roy MJ, Whitworth C, Diers E, Trainor N, et al. BAF complex vulnerabilities in cancer demonstrated via structure-based PROTAC design. *Nat Chem Biol.* 2019;15(7):672–680.
- [42] Kofink C, Trainor N, Mair B, Wöhrle S, Wurm M, Mischerikow N, et al. A selective and orally bioavailable VHL-recruiting PROTAC achieves SMARCA2 degradation in vivo. *Nat Commun.* 2022;13(1):5969.
- [43] Hulse M, Wang M, Xu C, Carter J, Agarwal A, Lu L, et al. PRT3789 is a first-in-human SMARCA2-selective degrader that induces synthetic lethality in SMARCA4-mutated cancers. *Cancer Res.* 2026;86(1):213–235.
- [44] van der Lelij P, Lieb S, Jude J, Wutz G, Santos CP, Falkenberg K, et al. Synthetic lethality between the cohesin subunits STAG1 and STAG2 in diverse cancer contexts. *eLife.* 2017;6:e26980.
- [45] Tsherniak A, Vazquez F, Montgomery PG, Weir BA, Kryukov G, Cowley GS, et al. Defining a cancer dependency map. *Cell.* 2017;170(3):564–576.e16.
- [46] Behan FM, Iorio F, Picco G, Gonçalves E, Beaver CM, Migliardi G, et al. Prioritization of cancer therapeutic targets using CRISPR–Cas9 screens. *Nature.* 2019;568(7753):511–516.
- [47] Gonatopoulos-Pournatzis T, Aregger M, Brown KR, Farhangmehr S, Braunschweig U, Ward HN, et al. Genetic interaction mapping and exon-resolution functional genomics with a hybrid Cas9–Cas12a platform. *Nat Biotechnol.* 2020;38(5):638–648.

- 1189 [48] Ito T, Young MJ, Li R, Jain S, Wernitznig A, Krill-Burger JM, et al. Paralog knockout profiling
1190 identifies DUSP4 and DUSP6 as a digenic dependence in MAPK pathway-driven cancers. *Nat*
1191 *Genet.* 2021;53(12):1664–1672.
- 1192 [49] Chou J, Esmaili Anvar N, Elghaish R, Chen J, Hart T. Z-scores outperform similar methods for
1193 analyzing CRISPR paralog synthetic lethality screens. *Genome Biol.* 2025;26(1):188.
- 1194 [50] Nijhawan D, Zack TI, Ren Y, Strickland MR, Lamothe R, Schumacher SE, et al. Cancer vulnera-
1195 bilities unveiled by genomic loss. *Cell.* 2012;150(4):842–854.
- 1196 [51] Muller FL, Colla S, Aquilanti E, Manzo VE, Genovese G, Lee J, et al. Passenger deletions generate
1197 therapeutic vulnerabilities in cancer. *Nature.* 2012;488(7411):337–342.
- 1198 [52] Cancer Dependency Map. DepMap 26Q1 Public Release [Dataset]. Broad Institute; 2026. [https:](https://depmap.org/portal/download/)
1199 [//depmap.org/portal/download/](https://depmap.org/portal/download/). Accessed 24 May 2026.
- 1200 [53] Meyers RM, Bryan JG, McFarland JM, Weir BA, Sizemore AE, Xu H, et al. Computational cor-
1201 rection of copy number effect improves specificity of CRISPR-Cas9 essentiality screens in cancer
1202 cells. *Nat Genet.* 2017;49(12):1779–1784.
- 1203 [54] Dempster JM, Boyle I, Vazquez F, Root DE, Boehm JS, Hahn WC, et al. Chronos: a cell population
1204 dynamics model of CRISPR experiments that improves inference of gene fitness effects. *Genome*
1205 *Biol.* 2021;22:343.
- 1206 [55] HUGO Gene Nomenclature Committee. HGNC complete gene set [Dataset]. European Bioinfor-
1207 matics Institute; 2026. <https://www.genenames.org/>. Accessed 24 May 2026.
- 1208 [56] Liu J, Lichtenberg T, Hoadley KA, Poisson LM, Lazar AJ, Cherniack AD, et al. An integrated
1209 TCGA pan-cancer clinical data resource to drive high-quality survival outcome analytics. *Cell.*
1210 2018;173(2):400–416.e11.
- 1211 [57] Benjamini Y, Hochberg Y. Controlling the false discovery rate: a practical and powerful approach
1212 to multiple testing. *J R Stat Soc B.* 1995;57(1):289–300.
- 1213 [58] Hollstein PE, Eichner LJ, Brun SN, Kamireddy A, Svensson RU, Vera LI, et al. The AMPK-related
1214 kinases SIK1 and SIK3 mediate key tumor-suppressive effects of LKB1 in NSCLC. *Cancer Discov.*
1215 2019;9(11):1606–1627.

- 1216 [59] Pedregosa F, Varoquaux G, Gramfort A, Michel V, Thirion B, Grisel O, et al. Scikit-learn: machine
1217 learning in Python. *J Mach Learn Res*. 2011;12:2825–2830.
- 1218 [60] Harris CR, Millman KJ, van der Walt SJ, Gommers R, Virtanen P, Cournapeau D, et al. Array
1219 programming with NumPy. *Nature*. 2020;585(7825):357–362.
- 1220 [61] Virtanen P, Gommers R, Oliphant TE, Haberland M, Reddy T, Cournapeau D, et al. SciPy 1.0:
1221 fundamental algorithms for scientific computing in Python. *Nat Methods*. 2020;17(3):261–272.
- 1222 [62] National Cancer Institute. Proteomic Data Commons [Dataset]. <https://pdc.cancer.gov/>. Ac-
1223 cessed 24 May 2026.
- 1224 [63] Goldman MJ, Craft B, Hastie M, Repečka K, McDade F, Kamath A, et al. Visualizing and inter-
1225 preting cancer genomics data via the Xena platform. *Nat Biotechnol*. 2020;38(6):675–678.